

Cross-validating an open-source Rust FEM solver against Palace on a superconducting transmon geometry

Robb Walters

July 14, 2026

Abstract

Palace, the open-source finite-element electromagnetics solver from the AWS Center for Quantum Computing, has become a reference tool for superconducting-qubit design, yet no independent quantitative validation of it exists in the literature. We supply one. GEODE-FEM, an independently implemented open-source Rust finite-element solver, reproduces Palace’s eigenmode solution on a real transmon-plus-readout-resonator geometry — generated by DeviceLayout.jl’s `SingleTransmon` example and meshed once, so both solvers consume the identical mesh and the identical lumped-element junction model — to within 0.03% across all physical modes, with the junction LC mode agreeing to 0.000%. Exact analytic tripwires (the LC anchor and the Josephson $1/\sqrt{2}$ inductance scaling) guard against coincidental agreement. On an 8-core CPU node, geode-fem’s serial direct-factorization shift-invert eigensolve is roughly $4\times$ more efficient per core than Palace’s MPI-distributed Krylov solve at this 133k-DOF size, while Palace’s parallelism wins $1.7\times$ at the whole-node level — an architecture trade-off, not a ranking. Two caveats are stated up front: the model contains no ~ 4 GHz qubit mode *by construction* (both solvers agree on its absence), and geode-fem’s un-projected Lanczos admits one disclosed spurious junction-localized mode that Palace’s divergence-free projection suppresses. The full fixture-to-comparison pipeline is scripted and committed.

1 Introduction

Three-dimensional finite-element electromagnetics is now a load-bearing step in superconducting-qubit design: mode frequencies, dispersive shifts, and participation ratios extracted from field simulation feed directly into device Hamiltonians (Koch et al., 2007; Mineev et al., 2021a; Nigg et al., 2012). The open-source solver of record for this workflow is Palace (Palace developers, AWS Center for Quantum Computing, 2023), an MFEM-based (Anderson et al., 2021; Andrej et al., 2024) parallel solver from the AWS Center for Quantum Computing, typically driven by geometry from DeviceLayout.jl (AWS Center for Quantum Computing, 2025; Pairs and Carson, 2025). This stack is widely used and well engineered — and, unusually for scientific software of its influence, has never been the subject of a peer-reviewed or preprint publication: Palace exists as a repository, a blog post (AWS Quantum Technologies Blog, 2023), and a workshop talk; the DeviceLayout.jl transmon workflow exists as a blog post and a conference talk (Pairs, 2025; Pairs and Carson, 2025). Independent, quantitative, third-party validation of the stack does not exist in the literature.

This paper supplies that validation, and does so under conditions designed to make agreement meaningful rather than merely reassuring. We compare Palace against GEODE-FEM (“geode-fem”), an open-source finite-element solver implemented independently in Rust — different language, different assembly path, different eigensolver, different linear algebra — on a real transmon-plus-readout-resonator geometry produced by DeviceLayout.jl’s own `SingleTransmon` example. The geometry is meshed *once*; both solvers consume the identical mesh, the identical first-order

discretization, and the identical lumped reactive-shunt junction model. Agreement claims are therefore about formulation and solver correctness, not meshing luck.

Our contributions are:

- **The first independent validation of the Palace/DeviceLayout.jl transmon stack.** On the identical mesh at matched order, geode-fem reproduces Palace’s physical eigenmodes to $\leq 0.03\%$, with the junction LC mode agreeing to 0.000% (Section 6).
- **A reproducible same-mesh cross-solver methodology.** An oracle-first benchmark protocol — exact analytic tripwires, inverse tests that must fail, participation-based mode identification, and a fully scripted fixture-to-comparison pipeline (Sections 5 and 11).
- **An honest architecture-level performance comparison.** On an 8-core node, a serial direct-factorization shift-invert eigensolve is $\sim 4\times$ more efficient per core than a distributed Krylov solve at 133k DOFs, while MPI parallelism wins $1.7\times$ at the whole-box level (Section 9); a GPU cell is scoped but pending (Section 10).
- **Honest-negative reporting as a first-class result.** The model’s lack of a ~ 4 GHz qubit mode is explained and confirmed by both solvers (Section 7); a spurious mode leaked by geode-fem’s un-projected Lanczos is disclosed, characterized, and tied to a known fix (Section 8).

Section 2 positions this work; Section 3 defines the benchmark problem; Section 4 contrasts the two solver architectures; the remaining sections present methodology, results, and reproducibility details.

2 Related Work

Electromagnetic simulation workflows for superconducting qubits. The closest prior work consists of two recent Palace-based workflow papers. Sommers et al. (2025) present SQDMetal, an open workflow that couples Qiskit Metal, Gmsh, and Palace, and cross-check Palace against the commercial solvers COMSOL and Ansys HFSS on qubit and resonator geometries, reporting agreement at roughly the 0.02–0.13% level. Ye et al. (2025) build a Palace-centered layout-to-Hamiltonian workflow and benchmark it against cryogenic measurement, with resonator frequencies reproduced to within 0.3%. Both efforts treat Palace as a component to be *wrapped and trusted*: validation runs against commercial tools or experiment, through independently constructed models and meshes, so residuals mix solver error with meshing and model-entry differences. Neither is a same-mesh, same-formulation comparison against an independently *implemented* solver. Our study eliminates meshing and model entry as confounders — identical MSH fixture, identical lumped-shunt junction, first-order Nédélec elements on both sides — so residuals isolate formulation and solver correctness; to our knowledge it is the first independent third-party validation of Palace itself. The broader open ecosystem for qubit electromagnetic design includes Qiskit Metal and its quasi-lumped analysis methods (Minev et al., 2021b) and the SQuADDs validated design database (Shanto et al., 2024); the physics framing for extracting qubit parameters from field modes rests on the transmon (Koch et al., 2007), black-box quantization (Nigg et al., 2012), and energy-participation-ratio (Minev et al., 2021a) literature.

Verification, validation, and cross-solver benchmarking. Code-to-code comparison on precisely specified benchmark problems is an established genre in computational electromagnetics: the TEAM (Testing Electromagnetic Analysis Methods) workshop tradition has run such comparisons

since 1986.¹ In the V&V vocabulary of Roache (1997) and Oberkampf and Roy (2010), the present study is verification-side evidence: our analytic tripwires are code-verification instruments in exactly Roache’s sense, and the same-mesh protocol is a code-to-code comparison in Oberkampf and Roy’s taxonomy. Multi-solver comparisons on shared electromagnetic problems have precedent in nano-optics, where Hoffmann et al. (2009) compared FEM, FDTD, and FIT codes on plasmonic antennas. We import that discipline to the superconducting-qubit domain and add two elements the genre often lacks: exact analytic tripwires that must move in known ways, and honest-negative reporting of what the model does *not* contain.

Open-source FEM/EM solver ecosystems. Mature open electromagnetics tools are normally accompanied by a citable methods paper: GetDP (Dular et al., 1998), DOLFIN/FEniCS (Logg and Wells, 2010), Gmsh (Geuzaine and Remacle, 2009), and scikit-rf (Arsenovic et al., 2022) all have one; so does Palace’s substrate library MFEM (Anderson et al., 2021; Andrej et al., 2024). Palace and DeviceLayout.jl are outliers — the *library* is published, the solver product built on it is not — which is precisely the gap this benchmark narrows from the outside, at a formality level (preprint) exceeding the original stack’s blog-and-talk record.

Numerics. geode-fem’s eigenpath is a shift-invert Lanczos (Lehoucq et al., 1998) over a real symmetric pencil, factorized with the faer sparse-LU (Kazdadi, 2026) and assembled on the Burn tensor framework (Tracel AI, 2024); the element family is first-order Nédélec (Nédélec, 1980). Palace’s eigenpath is a SLEPc-class Krylov method (Hernandez et al., 2005) with divergence-free projection and auxiliary-space (AMS) preconditioning (Hiptmair and Xu, 2007; Kolev and Vassilevski, 2009).

3 Benchmark Problem: Geometry, Mesh, and Junction Model

3.1 Geometry and mesh provenance

The device is the `SingleTransmon` example shipped with DeviceLayout.jl v1.15.0 (AWS Center for Quantum Computing, 2025): a transmon with readout resonator on a sapphire substrate, enclosed in a vacuum box. The sapphire is anisotropic, with permittivity tensor $\varepsilon = R \text{diag}(9.3, 9.3, 11.5) R^T$ rotated approximately 36.87° in-plane. The DeviceLayout.jl workflow emits seven named physical groups (metal, substrate, vacuum, junction lumped-element surface, two readout ports, exterior boundary) and drives Gmsh (Geuzaine and Remacle, 2009) to produce a tetrahedral mesh, exported as an MSH 4.1 fixture and pinned by SHA-256 (Section 11).

The mesh has 22,684 nodes and 133,314 tetrahedra, yielding 156,863 Nédélec edge unknowns, of which 133,108 interior degrees of freedom remain after eliminating perfect-electric-conductor (PEC) boundary edges. PEC is applied on all metal surfaces and the exterior box; the two readout ports are left open (lossless) in this first version, so no resistive element enters and the eigenproblem pencil stays real symmetric. Figure 1 shows the geometry with physical groups color-coded.

3.2 The junction model: a lumped reactive shunt

The Josephson junction is modeled identically in both solvers as a lumped reactive shunt — a parallel $L \parallel C$ element with $L = 14.860 \text{ nH}$ and $C = 5.5 \text{ fF}$ — attached to the four-triangle `lumped_element` surface group Γ . Following the uniform lumped-port construction, with port length ℓ and width w , the shunt contributes

$$K_{\text{port}} = \frac{\ell}{w L} S_\Gamma \quad \text{and} \quad M_{\text{port}} = \frac{C \ell}{w} S_\Gamma, \quad (1)$$

¹The TEAM founding documents are workshop reports and proceedings without resolvable DOIs; we describe the tradition in prose rather than cite an arbitrary secondary source.




Figure 1: Benchmark geometry: the DeviceLayout.jl v1.15.0 `SingleTransmon` transmon and readout resonator, with the seven named physical groups color-coded. Both solvers consume the identical mesh generated from this geometry (22,684 nodes; 133,314 tetrahedra).

Table 1: Solver architectures compared. The two implementations share the discretized problem definition (identical mesh, first-order Nédélec elements, identical lumped-shunt junction) and are otherwise independent.

	geode-fem	Palace
Language / substrate	Rust / Burn tensors (f64)	C++ / MFEM
Assembly	sparse full-tensor Nédélec	MFEM $H(\text{curl})$
Eigensolver	shift-invert Lanczos	SLEPc-class Krylov
Inner linear solve	faer sparse LU (direct)	AMS-preconditioned iterative
Spurious-mode control	none (post-hoc filter; §8)	divergence-free projection
Parallelism	single process	MPI-distributed

where S_Γ is the port-surface matrix assembled on Γ . The inductive term is frequency-independent and adds to the stiffness matrix; the capacitive term adds to the mass matrix. The eigenproblem is then the real symmetric generalized pencil

$$(K + K_{\text{port}})x = \omega^2 (M + M_{\text{port}})x. \quad (2)$$

In Palace, the same physics is entered as a reactive `LumpedPort` ($L \parallel C$, no resistance) on the same surface group; the readout ports are configured open (no $50\,\Omega$ termination), so Palace’s pencil is likewise real. This one-to-one correspondence of the junction treatment is what makes the comparison same-physics as well as same-mesh.

4 Solver Architectures Under Comparison

The two solvers share the mathematical problem (Eq. 2) and nothing else: no code, no language, no linear algebra, no eigensolver strategy. Table 1 summarizes the contrast.

geode-fem assembles the full-tensor Nédélec system sparsely in double precision, reduces to interior DOFs, and solves the pencil with a real shift-invert Lanczos iteration whose inner solves are direct sparse-LU factorizations. It runs in a single process with no intra-solve parallelism today. Palace assembles via MFEM, applies a divergence-free projection to confine the Krylov iteration to the physical subspace, preconditions with auxiliary-space methods (AMS) (Hiptmair and Xu, 2007; Kolev and Vassilevski, 2009), and distributes over MPI ranks. These choices have consequences for both correctness behavior (Section 8) and performance (Section 9); the point of the benchmark is to expose those consequences honestly rather than to declare a winner.

5 Validation Methodology

The benchmark protocol is oracle-first: every claim is anchored either to an exact analytic quantity or to a committed reference artifact, and the test suite is designed so that silent pipeline failures are caught by construction.

Same-mesh, same-physics protocol. The mesh is generated once, committed, and pinned by SHA-256. Both solvers consume it at matched (first) order with the identical junction model (Section 3.2). This removes meshing and model entry as confounders: any residual disagreement is attributable to formulation or solver implementation.

Analytic tripwires. Two exact quantities guard against coincidental or self-consistent-but-wrong agreement. First, the isolated junction LC resonance $f_{LC} = 1/(2\pi\sqrt{LC}) = 17.60$ GHz anchors the expected location of the junction mode. Second, a Josephson inductance-scaling tripwire: doubling L must move the junction mode by exactly $1/\sqrt{2}$. In our runs the junction mode moves from 17.49 to 12.37 GHz, a ratio of $0.7071 = 1/\sqrt{2}$ to the reported precision (Figure 3).

Inverse tests that must fail. Alongside tests asserting agreement, the suite contains perturbed configurations (such as the doubled- L case) that are *required to disagree* with the committed oracle. A pipeline that still “passes” under such a perturbation is not measuring anything; the inverse tests make that failure mode loud.

Participation-based mode identification. Modes are identified, not eyeballed. For each eigenvector x the junction participation

$$p = \frac{x^\top K_{\text{port}} x}{x^\top (K + K_{\text{port}}) x} \quad (3)$$

cleanly separates the junction LC mode ($p = 1.000$) from the resonator and box modes ($p \leq 0.0005$). Palace’s per-port energy-participation output provides the same discriminant on its side, and the two assignments agree.

Determinism check. The Palace oracle is not a one-off: rerunning Palace from the committed configuration reproduces its committed `eig.csv` bit-for-bit.

6 Results: Cross-Solver Agreement

Table 2 presents the headline result: on the identical 133k-DOF mesh at matched order, every physical mode agrees to $\leq 0.03\%$, and the junction LC mode — the mode the junction model exists to produce — agrees to 0.000% .

Three features of Table 2 deserve comment. First, the junction mode sits at 17.4901 GHz in both solvers, 0.6% below the isolated-circuit anchor $f_{LC} = 17.60$ GHz — the small shift reflects the junction’s embedding in the field problem, and its near-coincidence with the analytic value in *both* solvers is the tripwire working as intended. Second, the mode identification is unambiguous: junction participation is $p = 1.000$ for the 17.49 GHz mode and $p \leq 0.0005$ for every other mode, in both solvers’ participation outputs. Third, the agreement is between genuinely independent code paths — a direct-factorization Lanczos in Rust and a projected, preconditioned Krylov method in C++ — so a common implementation error is an unlikely explanation for the residual $\leq 0.03\%$, and the analytic tripwires exclude self-consistent-but-wrong agreement about the junction physics.

Table 2: Eigenmode agreement on the identical mesh (22,684 nodes / 133,314 tets / 156,863 Nédélec DOFs, 133,108 interior after PEC), first order in both solvers, junction as a lumped reactive shunt ($L = 14.860$ nH, $C = 5.5$ fF). geode-fem at commit 3174015; Palace at commit fba6a5b.

Mode	geode-fem (GHz)	Palace (GHz)	Δ
readout resonator	5.1528	5.1513	0.029%
box/cavity modes	15.4650 / 18.6927 / 20.6976 / 26.0809	same to $\leq 0.03\%$	$\leq 0.03\%$
junction LC	17.4901	17.4901	0.000%




Figure 2: Mode-frequency agreement: geode-fem versus Palace on the identical mesh, with per-mode relative deviations annotated. Data from the committed benchmark artifacts (benchmarks/transmon_eigen/results.toml and Palace’s committed eig.csv).

7 The Absent Qubit Mode: An Expected Negative

A reader familiar with the DeviceLayout.jl blog workflow (Peairs and Carson, 2025) will notice that its reported spectrum spans [4.14, 5.591] GHz — including a ~ 4 GHz qubit mode — while Table 2 contains no mode below 5.15 GHz. This is not a discrepancy to be minimized; it is a property of the model, and we state it plainly: **no ~ 4 GHz qubit mode exists in this junction model, by construction.**

The physical transmon frequency arises from the Josephson inductance resonating against the *pad (shunt) capacitance* of roughly 80–100 fF (Koch et al., 2007). The v1 lumped-shunt model of Section 3.2 gives the junction only its own 5.5 fF, so the junction’s resonance lands at 17.5 GHz, not near 4 GHz; the pad capacitance exists in the field solution but is not wired across the lumped element. Both solvers agree on the absence: Palace’s projected eigensolve, instructed to hunt at a shift of $\sigma = 4.5$ GHz, finds nothing below the 5.15 GHz resonator mode. The blog’s [4.14, 5.591] GHz spectrum is therefore *not reproduced by this model*, and no agreement with it is claimed.

The route to genuine qubit quantities from this benchmark is energy-participation-ratio (EPR) post-processing over the field modes (Minev et al., 2021a; Nigg et al., 2012), in which the junction’s participation in each electromagnetic mode — exactly the quantity in Eq. 3 — feeds the quantization. That analysis is planned future work (Phase C of this benchmark program) and is out of scope for the present cross-validation.

8 A Disclosed Spurious Mode

Honest cross-validation requires reporting where the solvers *disagree*, and there is exactly one such place. geode-fem’s Lanczos iteration operates on the raw pencil of Eq. 2 without a divergence-free




Figure 3: Josephson inductance-scaling tripwire: junction-mode frequency versus junction inductance L on log-log axes. Doubling L moves the junction mode from 17.49 to 12.37 GHz — ratio $0.7071 = 1/\sqrt{2}$ — exactly as the $f \propto 1/\sqrt{L}$ scaling requires.




Figure 4: Junction-participation spectra, geode-fem versus Palace, on the identical mesh. The physical junction LC mode ($p = 1.000$, 17.49 GHz) appears in both solvers; geode-fem’s un-projected Lanczos additionally leaks one spurious junction-localized mode near 3.45 GHz ($p = 0.994$) that Palace’s divergence-free projection suppresses.

projection. On this problem it leaks one spurious mode at approximately 3.45 GHz, below the physical band, localized on the junction surface with participation $p = 0.994$. Palace, whose eigensolve projects onto the divergence-free subspace, shows no such mode — its spectrum is clean down through the $\sigma = 4.5$ GHz hunt of Section 7.

We classify this mode as spurious — a gradient-nullspace-adjacent artifact of the un-projected iteration, not a physical resonance — on documented criteria: (i) it is absent from Palace’s projected spectrum on the identical mesh; (ii) it sits far below the analytic junction anchor $f_{LC} = 17.60$ GHz while carrying near-total junction participation, a combination inconsistent with the physical junction mode that appears, in both solvers, at 17.49 GHz; and (iii) it does not respond as a physical mode under the tripwire perturbations. The mode is excluded from Table 2 by these criteria, which are recorded alongside the committed benchmark configuration. Figure 4 contrasts the two solvers’ participation spectra, spurious mode included.

The fix is known and flagged as future work: a divergence-free or tree-cotree gauge projection on geode-fem’s eigenpath, mirroring the projection Palace applies. We prefer disclosing the artifact and its filtering criteria over silently post-selecting the spectrum: the disagreement is itself a finding about what the divergence-free projection buys in a production eigensolver.

Table 3: CPU-cell wall-clock and memory, full pipeline (mesh load + assembly + solve + output), `m6i.4xlarge`. Uncertainties are over $n = 3$ runs where shown. Palace at 16 ranks (hyperthreads) was refused by MPI process binding and is excluded — it is not a data point.

Solver	Parallelism	Wall (s)	Peak RSS
geode-fem @3174015	1 process (serial faer sparse-LU shift-invert Lanczos)	51.2 ± 0.4	3.1 GB
Palace @fba6a5b	4 MPI ranks	50.8	~ 0.7 GB/rank
Palace @fba6a5b	8 MPI ranks	30.6 ± 0.1	~ 0.5 GB/rank




Figure 5: CPU-cell wall-clock comparison (geode-fem single process versus Palace at 4 and 8 MPI ranks), with a per-core-efficiency inset. Full pipeline timings from Table 3.

9 Performance: CPU Cell

9.1 Setup

All CPU measurements ran on a single Amazon EC2 `m6i.4xlarge` instance (8 physical cores / 16 vCPUs, 64 GB RAM) in `us-west-2`.² Each measurement covers the *full pipeline* — mesh load, assembly, solve, and output — timed with `/usr/bin/time`; entries with uncertainties are the mean over $n = 3$ runs. geode-fem ran at commit 3174015, Palace at commit fba6a5b, both release builds.

9.2 Results

Table 3 supports two statements that must be read together. Per core, geode-fem’s serial direct-factorization eigensolve is roughly $4\times$ more efficient than Palace’s distributed Krylov–Schur/AMS solve on this 133k-DOF problem: one geode-fem core matches four Palace ranks almost exactly (51.2 versus 50.8s). At the whole-box level, Palace’s MPI parallelism wins $1.7\times$ (30.6 versus 51.2s), because geode-fem has no intra-solve parallelism today and leaves seven cores idle.

This is an architecture trade-off, not a ranking. At 133k DOFs the sparse LU factorization fits comfortably in memory (3.1 GB peak) and its flop efficiency dominates; a distributed iterative method carries communication and preconditioner-setup overheads that cannot amortize at this size, but distributes memory (~ 0.5 GB/rank at 8 ranks) and scales to problem sizes and rank counts where a serial direct factorization is not viable at all. The honest summary is that each architecture wins the regime it was designed for, and this benchmark documents where the crossover sits for a production transmon geometry at first order. Palace at 16 ranks on the hyperthreaded vCPUs was refused by MPI process binding; we exclude it rather than report a degraded configuration as if it

²On-demand cost approximately \$0.77/hour at the time of the runs; the entire CPU cell is reproducible for a few dollars.

were a measurement.

10 Performance: GPU Cell (Pending)

[TBD-GPU: This entire section awaits measured numbers. No GPU timing or correctness figure below is a result; every value is a placeholder.]

10.1 Scope and honest framing

The GPU cell is scoped but not yet run; we state its design and its limitations now so the eventual numbers cannot be over-read. geode-fem’s GPU path today accelerates the *driven* (deterministic frequency-sweep) solve — a matrix-free matvec with an on-device COCG iteration — and *not* the eigensolve that produced the headline comparison of Section 6. A GPU column for the eigenmode benchmark would therefore be misleading; the GPU cell instead measures the driven-solve pipeline on the same fixture. Additionally, the CUDA backend (burn-cuda 0.21) is f32-only — the underlying cubec1 runtime currently disables f64³ — so all GPU results are mixed-precision-qualified and will be reported with explicit correctness deltas against the f64 CPU path, not as drop-in equivalents.

10.2 Planned measurement

Hardware: a single Amazon EC2 `g6e.xlarge` instance (1× NVIDIA L40S, 46 GB; CUDA 13.2). Pending at the time of drafting:

- geode-fem CUDA-f32 correctness smokes (matrix-free matvec, on-device COCG, driven-solve wiring): [TBD-GPU: pass/fail + tolerance table].
- Driven-solve wall-clock, CPU versus GPU, same fixture: [TBD-GPU: timing table].
- Palace-GPU (libCEED/CUDA build): [TBD-GPU: explicitly future work if not run for the revision — will be stated as such, not silently dropped].

[TBD-GPU: Figure 5 (GPU-cell results) is deferred until measured data exists.]

11 Reproducibility

Every artifact behind every number in this paper is committed and scripted.⁴

- **Geometry and mesh.** Generated by DeviceLayout.jl v1.15.0’s `SingleTransmon` example (AWS Center for Quantum Computing, 2025); the fixture generator is committed, and the MSH 4.1 fixture is pinned by SHA-256 (5b3ff4c3...b33dd; full hash and provenance in the committed `transmon_smoke.provenance.txt`).
- **The MSH conversion gotcha.** MFEM’s MSH 4.1 reader rejects the DeviceLayout multi-entity-block file (“vertices indices are not unique”), so the Palace runs consume an MSH 2.2 conversion (`gmsh -save -format msh2`). The conversion is node-preserving (all 22,684 nodes unchanged, physical groups preserved) and physics-neutral: eigenvalues reproduce bit-for-bit across the conversion. We document this because it is exactly the kind of silent pipeline step a same-mesh claim must disclose.

³Upstream constraint of the burn/cubec1 stack at version 0.21 (Tracel AI, 2024); tracking-issue URL to be added at revision. TODO(operator).

⁴Repository URL and archival DOI for the geode-fem source at commit 3174015 to be added at submission. TODO(operator).

- **Solver configurations.** The Palace configuration and its provenance file are committed under `reference/fixtures/transmon_palace/`; geode-fem’s benchmark definitions live in `benchmarks/transmon_eigen/results.toml`, alongside the exact benchmark commands.
- **Oracles.** Palace’s raw outputs (`eig.csv`, per-port participation, run log) are committed; a rerun from the committed configuration reproduces `eig.csv` bit-for-bit.
- **Hardware disclosure.** CPU cell: EC2 `m6i.4xlarge` (8 physical cores / 16 vCPU, 64 GB), `us-west-2`. GPU cell (pending): EC2 `g6e.xlarge` (1× NVIDIA L40S 46 GB, CUDA 13.2).
- **Versions.** geode-fem commit 3174015; Palace commit `fba6a5b`; DeviceLayout.jl v1.15.0; Burn/burn-cuda 0.21.

12 Discussion

What the agreement does and does not establish. In V&V terms (Oberkampf and Roy, 2010; Roache, 1997), this study is code verification and code-to-code comparison, not validation against nature: two independent implementations agreeing to $\leq 0.03\%$ on the identical discrete problem, with analytic tripwires pinning the junction physics. It does not certify that either solver’s answer matches a measured device — that is what Ye et al. (2025)-style experimental benchmarks address, and the two evidence types are complementary. What same-mesh agreement uniquely rules out is the large class of errors that live in formulation, assembly, boundary-condition handling, and eigensolver machinery.

Threats to validity. A common-mode error — both solvers wrong in the same way — cannot be excluded by cross-comparison alone; the analytic anchors (f_{LC} , the $1/\sqrt{2}$ scaling) and the genuinely disjoint implementations make it unlikely for the junction physics specifically. Configuration-entry mismatch (comparing subtly different problems) is mitigated by committed configurations, the shared mesh file, and participation-based mode identification on both sides. Performance numbers are $n = 3$ full-pipeline measurements on one instance type; they support the architecture-level reading in Section 9, not microsecond-level claims.

Limitations. (i) One geometry, one mesh, first order only — higher-order and mesh-refinement convergence is future work (a $p = 2$ Palace configuration is already committed for that purpose). (ii) The lossless v1 model leaves readout ports open; resistive terminations (and hence complex-symmetric pencils and mode Q s) are out of scope. (iii) The spurious-mode filter of Section 8 currently leans on the committed oracle and documented criteria; the principled fix (divergence-free/tree-cotree projection) is identified but not yet implemented. (iv) The per-core efficiency result is specific to the 133k-DOF regime where a serial direct factorization is memory-feasible; it will not extrapolate to problem sizes that demand distribution. (v) The GPU cell is pending; no GPU claims are made in this version.

13 Conclusion

We presented the first independent, quantitative validation of the Palace and DeviceLayout.jl superconducting-transmon simulation stack: an independently implemented Rust FEM solver reproduces Palace’s eigenmodes on the identical mesh and identical junction model to $\leq 0.03\%$ across all physical modes (0.000% on the junction LC mode), with exact analytic tripwires guarding the junction physics. The benchmark’s methodology — same-mesh protocol, oracle-first tripwires, inverse tests that must fail, participation-based mode identification, and honest-negative reporting

of both the by-construction absent qubit mode and a disclosed spurious mode — is itself a reusable template for cross-solver validation in this domain. The CPU performance cell documents a clean architecture trade-off: $\sim 4\times$ per-core efficiency for the serial direct-factorization eigensolve against a $1.7\times$ whole-box win for MPI distribution at 133k DOFs. A mixed-precision GPU cell for the driven solve is scoped and pending. Both solvers emerge with their correctness confirmed on this problem class — which is precisely the outcome an ecosystem that increasingly trusts open solvers for quantum-hardware design should demand, and until now could not cite.

Acknowledgments. geode-fem’s development was AI-orchestrated (agent-driven implementation with human direction and review); this benchmark program and paper follow the same discipline, with every reported number traced to a committed, human-auditable artifact.

References

- Robert Anderson, Julian Andrej, Andrew Barker, Jamie Bramwell, Jean-Sylvain Camier, Jakub Cervený, Veselin Dobrev, Yohann Dudouit, Aaron Fisher, Tzanio Kolev, Will Pazner, Mark Stowell, Vladimir Tomov, Ido Akkerman, Johann Dahm, David Medina, and Stefano Zampini. MFEM: a modular finite element methods library. *Computers & Mathematics with Applications*, 81:42–74, 2021. doi: 10.1016/j.camwa.2020.06.009.
- Julian Andrej, Nabil Atallah, Jan-Phillip Bäcker, Jean-Sylvain Camier, Dylan Copeland, Veselin Dobrev, Yohann Dudouit, Tobias Duswald, Brendan Keith, Dohyun Kim, Tzanio Kolev, Boyan Lazarov, Ketan Mittal, Will Pazner, Socratis Petrides, Syun’ichi Shiraiwa, Mark Stowell, and Vladimir Tomov. High-performance finite elements with mfem. *The International Journal of High Performance Computing Applications*, 2024. doi: 10.1177/10943420241261981. URL <https://doi.org/10.1177/10943420241261981>.
- Alexander Arsenovic, Julien Hillairet, Jackson Anderson, Henrik Forsten, Vincent Ries, Michael Eller, Noah Sauber, Robert Weikle, William Barnhart, and Franz Forstmayr. scikit-rf: An open source python package for microwave network creation, analysis, and calibration [speaker’s corner]. *IEEE Microwave Magazine*, 23(1):98–105, 2022. doi: 10.1109/mmm.2021.3117139. URL <https://doi.org/10.1109/mmm.2021.3117139>.
- AWS Center for Quantum Computing. DeviceLayout.jl: computer-aided design of quantum integrated circuits. <https://github.com/aws-cqc/DeviceLayout.jl>, 2025. v1.15.0 used for geometry/mesh generation in this work. No journal or arXiv artifact exists as of 2026-07-14 (litsearch absence check).
- AWS Quantum Technologies Blog. AWS releases open-source software Palace for cloud-based electromagnetics simulations of quantum computing hardware. <https://aws.amazon.com/blogs/quantum-computing/aws-releases-open-source-software-palace-for-cloud-based-electromagnetics-simulations-of-quantum-computing-hardware/>, 2023.
- P. Dular, C. Geuzaine, F. Henrotte, and W. Legros. A general environment for the treatment of discrete problems and its application to the finite element method. *IEEE Transactions on Magnetics*, 34(5):3395–3398, 1998. doi: 10.1109/20.717799. URL <https://doi.org/10.1109/20.717799>.
- Christophe Geuzaine and Jean-François Remacle. Gmsh: a 3-D finite element mesh generator with built-in pre- and post-processing facilities. *International Journal for Numerical Methods in Engineering*, 79(11):1309–1331, 2009. doi: 10.1002/nme.2579.

- Vicente Hernandez, Jose E. Roman, and Vicente Vidal. SLEPc: a scalable and flexible toolkit for the solution of eigenvalue problems. *ACM Transactions on Mathematical Software*, 31(3):351–362, 2005. doi: 10.1145/1089014.1089019.
- Ralf Hiptmair and Jinchao Xu. Nodal auxiliary space preconditioning in $H(\text{curl})$ and $H(\text{div})$ spaces. *SIAM Journal on Numerical Analysis*, 45(6):2483–2509, 2007. doi: 10.1137/060660588.
- Johannes Hoffmann, Christian Hafner, Patrick Leidenberger, Jan Hesselbarth, and Sven Burger. Comparison of electromagnetic field solvers for the 3d analysis of plasmonic nano antennas, 2009. URL <https://arxiv.org/abs/0907.3570>.
- Sarah Quiñones El Kazdadi. faer: A linear algebra library for the rust programming language. *Journal of Open Source Software*, 11(123):6099, 2026. doi: 10.21105/joss.06099. URL <https://doi.org/10.21105/joss.06099>.
- Jens Koch, Terri M. Yu, Jay Gambetta, A. A. Houck, D. I. Schuster, J. Majer, Alexandre Blais, M. H. Devoret, S. M. Girvin, and R. J. Schoelkopf. Charge-insensitive qubit design derived from the Cooper pair box. *Physical Review A*, 76:042319, 2007. doi: 10.1103/PhysRevA.76.042319.
- Tzanio V. Kolev and Panayot S. Vassilevski. Parallel auxiliary space AMG for $H(\text{curl})$ problems. *Journal of Computational Mathematics*, 27(5):604–623, 2009. doi: 10.4208/jcm.2009.27.5.013. URL <https://doi.org/10.4208/jcm.2009.27.5.013>. Crossref records the online-publication year 2018; print year 2009 used here (litsearch cosmetic fix, see version-dir refs.bib header).
- Richard B. Lehoucq, Danny C. Sorensen, and Chao Yang. *ARPACK Users’ Guide: Solution of Large-Scale Eigenvalue Problems with Implicitly Restarted Arnoldi Methods*. SIAM, 1998. doi: 10.1137/1.9780898719628.
- Anders Logg and Garth N. Wells. Dolfin. *ACM Transactions on Mathematical Software*, 37(2):1–28, 2010. doi: 10.1145/1731022.1731030. URL <https://doi.org/10.1145/1731022.1731030>.
- Zlatko K. Mineev, Zaki Leghtas, Shantanu O. Mundhada, Lysander Christakis, Ioan M. Pop, and Michel H. Devoret. Energy-participation quantization of Josephson circuits. *npj Quantum Information*, 7:131, 2021a. doi: 10.1038/s41534-021-00461-8.
- Zlatko K. Mineev, Thomas G. McConkey, Maika Takita, Antonio D. Corcoles, and Jay M. Gambetta. Circuit quantum electrodynamics (cqed) with modular quasi-lumped models, 2021b. URL <https://arxiv.org/abs/2103.10344>.
- Jean-Claude Nédélec. Mixed finite elements in $\mathbb{R}^3$. *Numerische Mathematik*, 35(3):315–341, 1980. doi: 10.1007/BF01396415.
- Simon E. Nigg, Hanhee Paik, Brian Vlastakis, Gerhard Kirchmair, S. Shankar, Luigi Frunzio, M. H. Devoret, R. J. Schoelkopf, and S. M. Girvin. Black-box superconducting circuit quantization. *Physical Review Letters*, 108:240502, 2012. doi: 10.1103/PhysRevLett.108.240502.
- William L. Oberkampf and Christopher J. Roy. Verification and validation in scientific computing. 2010. doi: 10.1017/cbo9780511760396. URL <https://doi.org/10.1017/cbo9780511760396>. Cambridge University Press.
- Palace developers, AWS Center for Quantum Computing. Palace: PArallel LARge-scale Computational Electromagnetics. <https://github.com/aws-labs/palace>, 2023. Open-source 3D

- finite-element solver; v0.13-class, git fba6a5b used in this work. No associated journal article exists as of 2026-07 (absence re-confirmed by litsearch, 2026-07-14).
- Greg Peairs. Schematic-driven design of a quantum processor with DeviceLayout.jl. JuliaCon Global 2025, <https://pretalx.com/juliacon-2025/talk/7EYFPN/>, 2025. Recording: <https://www.youtube.com/watch?v=-LohBVtFdwU>.
- Greg Peairs and Hugh Carson. Design quantum integrated circuits with open-source software DeviceLayout.jl from AWS. <https://aws.amazon.com/blogs/quantum-computing/design-quantum-integrated-circuits-with-open-source-software-devicelayout-jl-from-aws/>, 2025. AWS Quantum Technologies Blog; source of the transmon+resonator geometry and the reference workflow reproduced here.
- P. J. Roache. Quantification of uncertainty in computational fluid dynamics. *Annual Review of Fluid Mechanics*, 29(1):123–160, 1997. doi: 10.1146/annurev.fluid.29.1.123. URL <https://doi.org/10.1146/annurev.fluid.29.1.123>.
- Sadman Shanto, Andre Kuo, Clark Miyamoto, Haimeng Zhang, Vivek Maurya, Evangelos Vlachos, Malida Hecht, Chung Wa Shum, and Eli Levenson-Falk. Squadds: A validated design database and simulation workflow for superconducting qubit design. *Quantum*, 8:1465, 2024. doi: 10.22331/q-2024-09-09-1465. URL <https://doi.org/10.22331/q-2024-09-09-1465>.
- David Sommers, Zach Degnan, Divita Gautam, Yi-Hsun Chen, Chun-Ching Chiu, Arkady Fedorov, and Prasanna Pakkiam. Open-source highly parallel electromagnetic simulations for superconducting circuits, 2025. URL <https://arxiv.org/abs/2511.01220>.
- Tracel AI. Burn: a comprehensive deep-learning and tensor-computation framework in Rust. <https://github.com/tracel-ai/burn>, 2024. v0.21.0 used. No journal/arXiv paper exists (litsearch check 2026-07-14): the repo CITATION.cff lists authors Simard, Fortier-Dubois, Tadjibaev, Lagrange + Burn Framework Contributors, with no DOI and no preferred-citation paper — cited as software.
- Jiale Ye, Jiaheng Wang, and Yu-xi Liu. Electromagnetic feature extraction in superconducting quantum circuits: An open-source finite-element workflow using palace, 2025. URL <https://arxiv.org/abs/2511.09041>.