Differentiable transmon design with a tensor-native finite-element electromagnetics solver: gradient-based optimization of charging energy, cross-validated against Palace

Robb Walters* Crutcher Dunnavant†

October 11, 2026

Abstract

Superconducting-qubit electromagnetic design is an iterative guess-and-check loop: the workflow's production field solvers expose no derivatives of design figures-of-merit, so state-of-practice optimizers iterate time-consuming electromagnetic simulations and inject parameter updates from separate analytic physics models (Eriksson et al., 2025). We demonstrate the missing capability — *solver-derived* design gradients from a 3D finite-element electromagnetics solve — and use it for gradient-based optimization of a transmon's charging energy E_C. GEODE-FEM is a Rust solver built on the Burn tensor framework; a discrete-adjoint layer bridges the one stage its autodiff tape cannot reach (the sparse factorization), yielding a validated 2×2 sensitivity matrix — material ($\partial/\partial\varepsilon$) and geometry ($\partial/\partial X$) derivatives on both the electrostatic and the driven $H(\mathrm{curl})$ operators — each finite-difference-validated with mutation tests (relative errors 10^{-9}–10^{-5}). Because the capacitance $C = \phi^{\mathsf{T}} K \phi$ is variationally stationary, its shape derivative collapses to a Hellmann–Feynman-like explicit term; composing it with the capacitance-to-E_C chain, Newton iteration drives a capacitor geometry to a 0.2156 GHz target E_C, confirmed by an independent fresh solve to 1.4×10^{-15}. On the real DeviceLayout transmon, a finite-difference-validated quadratic-element shape gradient on an adaptively refined mesh drives the island capacitance C_Σ to the 89.9 fF design anchor in three accepted steps. That C_Σ is the scalar-permittivity electrostatic value, with the junction's 5.5 fF self-capacitance excluded. Its gradient is a hand-derived equilibrium-energy shape derivative, solved by conjugate gradients with an algebraic-multigrid preconditioner in native f64; it does not run on the Burn tape and uses no factorization. Seven held-out refinement rounds then bound the converged capacitance of that design rigorously from above (below 89.02 fF) and estimate it by extrapolation at about 88.8–88.9 fF; we do not certify a continuum value. A second design, with the target offset by that measured bias, converges below 90.08 fF, at most 0.20% above the anchor. Its extrapolated band of 89.84–89.96 fF contains the anchor by construction, because its refinement series is the first one rescaled; at slower assumed convergence rates it lands 0.25–0.57% below. Only the upper side is proven, and a two-sided bound is future work. The convergence study is a result in its own right. Under uniform refinement the scalar-permittivity capacitance of the unmodified device falls from 137.7 fF on the committed 133k-tetrahedron mesh to below 83.02 fF (rigorous; estimated at 74.14–79.62 fF). The gap to the ∼90 fF anchor that earlier versions called a model difference was therefore discretization error, and the converged capacitance of the committed geometry lies below the anchor, not above it. How far below depends on the permittivity model: at least 7.6% with scalar permittivity, and at least 0.4% with the sapphire tensor that the lumped-element chain uses (converged below 89.50 fF). Earlier coarse-mesh optimizations lowered

*These authors contributed equally.
†These authors contributed equally.

C_Σ to reach the anchor; they moved the wrong way and are kept as a cautionary precursor. Eigenmode/EPR derivatives remain roadmap. The correctness credential: from the identical input mesh, GEODE-FEM reproduces Palace's six eigenmodes above 4.5 GHz to $\leq 0.033\%$ (worst case 0.032%), and the transmon qubit mode at 3.45 GHz to 5×10^{-5} %. Earlier versions of this paper called that mode spurious and the qubit mode absent; a Palace run with its search target below the mode corrects both statements. The contribution is a capability, not a speed race — on raw performance the measured record shows no corner where GEODE-FEM beats Palace at scale, and we report it honestly. Every result traces to a committed artifact; statements that rest on a scratch probe, on a reading of Palace's source, or on inference are labelled as such.

1 Introduction

The hard problem in superconducting-qubit electromagnetic design is the *inverse map*: given target Hamiltonian parameters — a charging energy E_C, an anharmonicity α, coupling capacitances — produce a physical layout. Today that map is walked by hand: parameterize a geometry, run a 3D electromagnetic simulation, extract the parameters, compare, adjust, repeat. The state-of-practice automated workflow makes the loop's structure explicit: QDesignOptimizer (Eriksson et al., 2025) iterates time-consuming electromagnetic simulations of HFSS-class solvers and guides the parameter updates with *separate*, user-defined analytic physics models — precisely because the field solver at the center of the loop produces observables but no derivatives of those observables with respect to the design. The gradient the optimizer actually wants, $\partial(\text{figure of merit})/\partial(\text{geometry, materials})$ through the field solve itself, is structurally unavailable from the production solvers — commercial (HFSS, COMSOL) and open-source (Palace (Palace developers, AWS Center for Quantum Computing, 2023)) alike: their assembly and factorization pipelines are hand-written kernels with no derivative path, and none exposes an adjoint of its solve.

The surrounding literature brackets this gap without closing it. At the *lumped-circuit* level, differentiation is a solved problem: SQcircuit (Rajabzadeh et al., 2023) analyzes arbitrary superconducting circuits, and its gradient-based follow-up (Rajabzadeh et al., 2024) computes eigenvalue and eigenvector gradients of the circuit Hamiltonian via PyTorch autodiff, optimizing qubit designs in circuit-parameter space. But the map from geometry and materials *to* those lumped parameters still runs through non-differentiable 3D electromagnetic simulation. At the *field* level, differentiable electromagnetics is owned by photonics: adjoint and autodiff methods on FDTD and integral-equation solvers are a mature inverse-design industry (Hammond et al., 2022; Hughes et al., 2019; Mahlau et al., 2026; Molesky et al., 2018; Ponomareva et al., 2025; Schubert et al., 2022), and differentiable FEM at scale has been proven outside electromagnetics by JAX-FEM in solid mechanics (Xue et al., 2023). What none of these supplies is the regime superconducting-qubit (and, more broadly, RF) design lives in: *frequency-domain finite elements on body-fitted tetrahedra with anisotropic materials* — distributed, field-level EM gradients from the same FEM solve that produces the design observables. That wedge is, to our knowledge, uncontested; this paper occupies it.

GEODE-FEM ("geode-fem") is positioned to close the gap by construction rather than retrofit. It is a full-wave 3D $H(\text{curl})$/Nédélec (plus scalar electrostatic) finite-element solver written in Rust on the Burn tensor framework (Tracel AI, 2024): element assembly is batched $[n_{\text{elem}}, 6, 6]$ tensor algebra on an autodiff-capable substrate, so the derivative of assembly with respect to materials and node coordinates is available where a classical solver has only hand-written kernels. One stage breaks the tape: the sparse direct factorization (faer (Kazdadi, 2026)) at the heart of every solve. A discrete-adjoint layer closes exactly that gap — one extra (transpose-)solve, reusing the forward factorization, restores the full gradient (Section 8). The resulting capability is general — shape and material sensitivities for RF, photonic, and electrostatic device design — and the transmon is our demonstration vehicle because its unmet need is the clearest documented in the literature, not

because the method is qubit-specific.

A design-gradient claim from a new solver is only as credible as the solver's forward answers, so the paper carries its correctness credential up front: on a real transmon-plus-readout-resonator geometry meshed once and given as the same input file to both codes, geode-fem reproduces Palace — the MFEM-based (Anderson et al., 2021; Andrej et al., 2024) reference solver for superconducting-qubit design — to within 0.033% across the six eigenmodes Palace returns above 4.5 GHz (worst case 0.032%), and to 5×10^{-5} % on the transmon qubit mode at 3.45 GHz, with analytic scaling tripwires guarding against coincidental agreement (Section 5). Palace refines the mesh around the metal sheets internally, so the two discrete problems are not identical (Section 4.1). This parity is *a credential, not the contribution* — matching Palace is the price of entry — and it is, incidentally, the first independent third-party validation of the Palace/DeviceLayout.jl transmon stack, which has no publication of its own.

Our contributions are:

- **A validated 2×2 discrete-adjoint sensitivity matrix for FEM electromagnetics.** Material ($\partial/\partial\varepsilon$) and geometry ($\partial/\partial X$, exact $\partial K/\partial X$ via forward-mode duals through the element kernels) derivatives, on both the scalar electrostatic and the complex driven $H(\mathrm{curl})$ Nédélec operators — each one forward plus one adjoint solve sharing a single factorization, each validated against full-pipeline central finite differences with mutation tests that reject sign, conjugation, and dropped-term errors (Section 8).

- **A differentiable capacitance$\to E_C$ chain with a variational-stationarity simplification.** Because $C = \phi^{\mathsf{T}} K \phi$ is stationary in the solved potential, the implicit term vanishes and the shape derivative is a pure explicit-geometry (Hellmann–Feynman-like) term; validated against the analytic parallel-plate law and central finite differences (Section 8.2).

- **A refinement-validated real-device optimization, with its error budget stated.** An end-to-end Newton loop hits a 0.2156 GHz target E_C on an analytic fixture and is confirmed by an independent fresh solve at 1.4×10^{-15} (the parametrization is affine, so one-step convergence proves the loop, not a hard optimization). On the real device, a three-parameter optimization on an adaptively refined 448k-tetrahedron mesh with quadratic elements reaches the 89.9 fF anchor in three accepted steps, with gradients finite-difference-validated at 3.87×10^{-7} or better. That gradient is a hand-derived equilibrium-energy shape derivative, solved by conjugate gradients with an algebraic-multigrid preconditioner in native f64. It uses the stationarity identity of the previous item, and it runs outside the Burn tape and without the adjoint layer's factorization. Seven held-out refinement rounds give a rigorous upper bound on the converged capacitance of the design (below 89.02 fF) and an extrapolated estimate (about 88.8–88.9 fF). A resolution-corrected second design converges below 90.08 fF. No continuum value is certified and no lower bound is proven (Section 9.4).

- **A measured mesh-convergence correction.** In the scalar-permittivity model, the committed 133k-tetrahedron mesh overstates C_Σ by at least 54.69 fF, and the converged unmodified device lies below 83.02 fF. With the sapphire tensor that the lumped-element chain uses, the overstatement is at least 47.18 fF and the converged value lies below 89.50 fF. Both bounds are under the ∼90 fF anchor, where earlier versions placed the device above it. Their "model difference" diagnosis is withdrawn, and the coarse-mesh optimizations that lowered C_Σ toward the anchor are re-read as a cautionary precursor (Sections 7 and 9.3).

- **The correctness credential.** Same-input-mesh, same-physics, same-junction-model cross-validation against Palace at $\leq 0.033\%$ across the six modes above 4.5 GHz and 5×10^{-5} %

on the qubit mode, with oracle-first methodology and honest-physics reporting — including the correction of two claims made in earlier versions, that the qubit mode was absent and that the 3.45 GHz mode was spurious (Sections 5 and 6) — doubling as the first independent validation of the Palace/DeviceLayout.jl stack.

- **Mode-verified performance context.** Same-box CPU measurements disclose mode selection, field-output and error-estimator costs; they replace the earlier unsupported one-core-versus-eight-rank efficiency claim. CPU-f64 AMS reduces driven-solve time relative to Jacobi, but its iteration and setup costs still grow with refinement (Section 11). Eigenmode/EPR derivatives remain roadmap (Section 10).

Section 2 positions the work; Section 3 presents the substrate and where its tape breaks; Sections 4–6 establish the benchmark problem and the correctness credential; Sections 7–9 are the contribution — the forward LOM chain with its mesh-convergence study, the adjoint layer, and the optimization tiers; Section 10 is the eigenmode roadmap; Sections 11–13 report performance context, reproducibility, and discussion.

2 Related Work

Superconducting-qubit design optimization. The design loop this paper attacks is documented across the qubit-design software stack. QDesignOptimizer (Eriksson et al., 2025) is the state of practice: physics-guided multi-parameter optimization that wraps non-differentiable EM simulation (HFSS with pyEPR/Qiskit-Metal-class post-processing) and steers it with analytic models, because no gradient is available from the solver. The SQcircuit line (Rajabzadeh et al., 2023, 2024) differentiates at the lumped-circuit level — including eigenvalue/eigenvector gradients of the circuit Hamiltonian via PyTorch — and optimizes in circuit-parameter space; the geometry-to-parameters map remains simulation-bound. The SQuADDS database (Shanto et al., 2024) attacks the same loop by lookup: pre-simulated, experimentally-validated design points. Qiskit Metal's quasi-lumped methods (Minev et al., 2021b) and the EPR/black-box quantization physics (Koch et al., 2007; Minev et al., 2021a; Nigg et al., 2012) frame the parameter-extraction layer this paper's LOM chain implements. None of these obtains gradients from the field solve itself; that is the gap Section 8 closes.

Differentiable electromagnetics — owned by photonics. Adjoint/autodiff EM is a mature inverse-design discipline in nanophotonics (Molesky et al., 2018), expressed through structured grids and integral methods: autograd FDFD/FDTD in ceviche (Hughes et al., 2019), foundry-constrained JAX inverse design (Schubert et al., 2022), the Meep adjoint pipeline (Hammond et al., 2022), the JAX-native FDTDX framework (Mahlau et al., 2026), GPU-accelerated autodiff integral-method scattering in TorchGDM (Ponomareva et al., 2025), and differentiable RCWA (Kim et al., 2024). These codes accept rectilinear grids or Green's-function discretizations; the superconducting-qubit regime — body-fitted tetrahedra, anisotropic sapphire, $H(\mathrm{curl})$ conformity, frequency domain — is exactly where that literature stops. Differentiable FEM has been proven outside EM: JAX-FEM (Xue et al., 2023) (GPU-accelerated, autodiff-native continuum mechanics) and PyTorch-FEA (Liang et al., 2023) in biomechanics; concurrent work on general differentiable Galerkin assembly (TensorGalerkin (Wen et al., 2026), with the torch-sla adjoint sparse-solver library (Chi and Wen, 2026)) independently confirms the substrate-level thesis while reporting nodal magnetostatics rather than full-wave $H(\mathrm{curl})$ EM. To our knowledge, no prior work delivers field-level design gradients from a frequency-domain 3D FEM electromagnetics solve of the class RF/superconducting design uses.

FEM kernels by code generation, and tensor compilers. Generating FEM kernels from high-level descriptions is a two-decade tradition — FFC (Kirby and Logg, 2006), UFL (Alnæs et al., 2014), Firedrake (Rathgeber et al., 2017), TSFC (Homolya et al., 2018) — whose modern GPU expression is libCEED (Brown et al., 2021) and MFEM's partial-assembly path (Andrej et al., 2024). Palace itself runs on that domain-specific stack. geode-fem's substrate bet is categorically different: outsource kernels, scheduling, and backend portability to a *general-purpose ML tensor stack* (the Halide/TVM compiler lineage (Chen et al., 2018; Ragan-Kelley et al., 2013), reaching scientific computing via ML frameworks), and phrase assembly and matrix-free application as batched tensor algebra. For this paper the decisive property of that choice is not portability but *differentiability*: an autodiff-capable substrate makes the assembly derivative available by construction, which is what the adjoint layer of Section 8 composes with. Adjacent genera — simulation DSLs (Holl et al., 2020; Hu et al., 2019), neural-ansatz eigensolvers (Jiang et al., 2024), hand-written CUDA FEM (Meng et al., 2014) — trade away either the general-purpose stack or the discretization itself.

Validation and cross-solver benchmarking. Code-to-code comparison on precisely specified problems is an established genre in computational electromagnetics (the TEAM workshop tradition since 1986[1]; nano-optics solver comparisons (Hoffmann et al., 2009)); in V&V vocabulary (Oberkampf and Roy, 2010; Roache, 1997) our credential is verification-side, code-to-code evidence. On the qubit-EM axis the closest works are SQDMetal (Sommers et al., 2025) (Palace cross-checked against COMSOL/HFSS at the 0.02–0.13% level, independently constructed models) and the Palace-centered workflow of Ye et al. (2025) (validated against cryogenic measurement at ∼0.3%). Both wrap and trust Palace; neither is a same-mesh comparison against an independently *implemented* solver. Our protocol removes mesh generation and model entry as confounders — identical MSH fixture, identical lumped-shunt junction, first-order Nédélec on both sides — so residuals reflect formulation, solver and Palace's internal refinement of the metal sheets (Section 4.1), not independently built models. Mature open EM tools normally carry a methods paper (GetDP (Dular et al., 1998), FEniCS/DOLFIN (Logg and Wells, 2010), Gmsh (Geuzaine and Remacle, 2009), scikit-rf (Arsenovic et al., 2022), MFEM (Anderson et al., 2021)); Palace and DeviceLayout.jl are unpublished outliers (AWS Center for Quantum Computing, 2025; AWS Quantum Technologies Blog, 2023; Palace developers, AWS Center for Quantum Computing, 2023; Peairs, 2025; Peairs and Carson, 2025), which is the gap the credential narrows from the outside. geode-fem's numerics: first-order Nédélec elements (Nédélec, 1980), shift-invert Lanczos (Lehoucq et al., 1998) over a real symmetric pencil factorized with faer (Kazdadi, 2026); Palace: SLEPc-class Krylov (Hernandez et al., 2005) with divergence-free projection. The timed Palace build resolves its default linear solver to SuperLU under GMRES; explicit AMS runs are reported separately (Hiptmair and Xu, 2007; Kolev and Vassilevski, 2009); the gauge-fixing literature behind Section 6 is the tree–cotree tradition (Albanese and Rubinacci, 1988; Manges and Cendes, 1995).

3 GEODE-FEM: A Tensor-Native Substrate, and Where the Tape Breaks

geode-fem is written in Rust on the Burn tensor framework (Tracel AI, 2024), whose cubecl JIT compiles tensor kernels for CPU, CUDA, and WebGPU backends from one source. Every throughput-critical FEM stage is phrased as the batched dense algebra the stack already optimizes.

Batched element-local assembly. For first-order Nédélec tetrahedra the element curl–curl and mass contributions are 6×6 matrices depending on per-element geometry factors and material tensors

[1]The TEAM founding documents are workshop reports without resolvable DOIs; we describe the tradition in prose.

(including the rotated anisotropic sapphire permittivity of Section 4.1); the scalar electrostatic operator is the analogous P1 kernel. geode-fem computes them for all elements at once as batched $[n_{\text{elem}}, 6, 6]$ tensor operations — one kernel launch per assembly stage rather than an element loop — then either scatters into a global sparse matrix (the assembled f64 CPU path) or retains the batch for matrix-free application.

Matrix-free application and on-device Krylov. The matrix-free operator apply is a three-stage tensor pipeline — gather global DOFs to element-local vectors, batched matmul against the stored element operators, scatter-add back — and the driven frequency-sweep solve runs a conjugate orthogonal conjugate gradient (COCG) iteration entirely on the device with a constant number of scalar host–device synchronizations per iteration. The backend (ndarray on CPU, CUDA on GPU) is a type parameter, not a code path. Section 11 reports honestly how far below the GPU-amortization threshold the current benchmark sizes sit.

Honest constraints. Two boundaries are constraints, not footnotes. The CUDA backend (burn-cuda 0.21) is f32-only today — the underlying cubecl runtime currently disables f64[2] — so every GPU result is mixed-precision-qualified. And the eigensolve that produces the credential comparison is factorization-bound: its inner solves are faer sparse-LU factorizations (Kazdadi, 2026), a classical serial CPU component outside the tensor-compiler story.

Where the tape breaks — and why an adjoint layer, not autodiff-through-the-solver. The substrate choice has a consequence the rest of this paper builds on: Burn's automatic differentiation reaches everything *up to* the linear solve — material and node-coordinate dependence of the assembled operators is differentiable by construction — but the sparse factorization is an external, opaque kernel that breaks the tape. Differentiating *through* a factorization by unrolling is neither practical nor necessary: the classical discrete-adjoint identity turns the missing derivative into one additional linear solve against the *same* factorization (Section 8). The division of labor is deliberate — autodiff where the substrate provides it (element kernels, via forward-mode duals), adjoint algebra where it does not (the solve) — and it is what makes the adjoint-layer gradients of Section 8 and the coarse-mesh tiers of Section 9 *solver-derived* in the sense the design loop of Section 1 lacks: the same assembly and the same factorization that produce the observable produce its gradient. The refined real-device optimization of Section 9.4 is solver-derived in a narrower sense. Its gradient is evaluated from the same solved field that gives the capacitance, but with hand-derived quadratic-element integrals and an iterative solve in native f64; it uses neither the Burn tape nor a factorization.

4 Benchmark Problem: Geometry, Mesh, and Junction Model

4.1 Geometry and mesh provenance

The device is the `SingleTransmon` example shipped with DeviceLayout.jl v1.15.0 (AWS Center for Quantum Computing, 2025): a transmon with readout resonator on a sapphire substrate, enclosed in a vacuum box. The sapphire is anisotropic, with permittivity tensor $\varepsilon = R\,\mathrm{diag}(9.3,\, 9.3,\, 11.5)\,R^{\mathsf{T}}$ rotated approximately 36.87° in-plane. The DeviceLayout.jl workflow emits seven named physical groups (metal, substrate, vacuum, junction lumped-element surface, two readout ports, exterior boundary) and drives Gmsh (Geuzaine and Remacle, 2009) to produce a tetrahedral mesh, exported as an MSH 4.1 fixture and pinned by SHA-256 (Section 12).

[2]Upstream constraint of the burn/cubecl stack at version 0.21 (Tracel AI, 2024); tracking-issue URL to be added at revision. TODO(operator).

The mesh has 22,684 nodes and 133,314 tetrahedra, yielding 156,863 Nédélec edge unknowns, of which 133,108 interior degrees of freedom remain after eliminating perfect-electric-conductor (PEC) boundary edges. PEC is applied on all metal surfaces and the exterior box; the two readout ports are left open (lossless), so no resistive element enters and the eigenproblem pencil stays real symmetric. These counts are geode-fem's. Palace reads the same file but does not solve the same discrete problem: its run logs record 4,913 elements added by local bisection around under-resolved interior boundaries and 6,014 duplicated vertices where it cracks the metal sheets, and it assembles 183,891 first-order Nédélec unknowns against geode-fem's 156,863 (about 17% more). "Same mesh" in this paper therefore means the same *input* mesh. Earlier versions said the two solvers consumed an identical mesh and solved an identical discrete problem; that was wrong, and the timing comparisons of Section 11.1 carry the difference as a fairness caveat. This mesh is also coarse for the electrostatic quantities: one uniform refinement lowers C_Σ by 26.3% (Section 7). Figure 1 shows the geometry with physical groups color-coded. Every measurement in this paper starts from this one committed mesh. The eigenmode credential and the coarse-mesh electrostatic chain run on it directly; the scale study of Section 11 uses a uniformly refined variant; and the convergence study and refined optimization of Sections 7 and 9.4 use uniform and adaptive refinements of it that leave its polygonal surfaces unchanged.

4.2 The junction model: a lumped reactive shunt

The Josephson junction is modeled identically in both solvers as a lumped reactive shunt — a parallel $L \parallel C$ element with $L = 14.860\,\mathrm{nH}$ and $C = 5.5\,\mathrm{fF}$ — attached to the four-triangle `lumped_element` surface group Γ. Following the uniform lumped-port construction, with port length ℓ and width w, the shunt contributes

$$K_{\mathrm{port}} = \frac{\ell}{w\,L}\,S_\Gamma \qquad \text{and} \qquad M_{\mathrm{port}} = \frac{C\,\ell}{w}\,S_\Gamma, \tag{1}$$

where S_Γ is the port-surface matrix assembled on Γ. The eigenproblem is then the real symmetric generalized pencil

$$(K + K_{\mathrm{port}})\,x \;=\; \omega^2\,(M + M_{\mathrm{port}})\,x. \tag{2}$$

In Palace, the same physics is entered as a reactive `LumpedPort` ($L \parallel C$, no resistance) on the same surface group; the readout ports are configured open, so Palace's pencil is likewise real. Palace's source confirms the correspondence: its `LumpedPort` adds the same scalar surface-mass terms with the same per-square scalings (`docs/research/transmon-lumped-port-formulation.md`, §1). This one-to-one correspondence is what makes the comparison same-physics as well as same-input-mesh. It also means that any property of the distributed surface term in Eq. 1 is shared by both solvers (Section 6.3).

5 The Correctness Credential: Same-Input-Mesh Cross-Validation Against Palace

A solver claiming design gradients must first prove its forward answers; this section is that proof. Palace's own agreement with commercial tools is established (Sommers et al., 2025), so reproducing Palace is the bar a new implementation must clear before anything it computes *about* its solutions (such as their derivatives) is worth discussing.

Protocol. The benchmark is oracle-first. The mesh is generated once, committed, pinned by SHA-256, and read by both solvers at matched (first) order with the identical junction model; residual disagreement is therefore attributable to formulation, implementation or Palace's internal refinement (Section 4.1), not to independent meshing. Two analytic tripwires guard against coincidental or

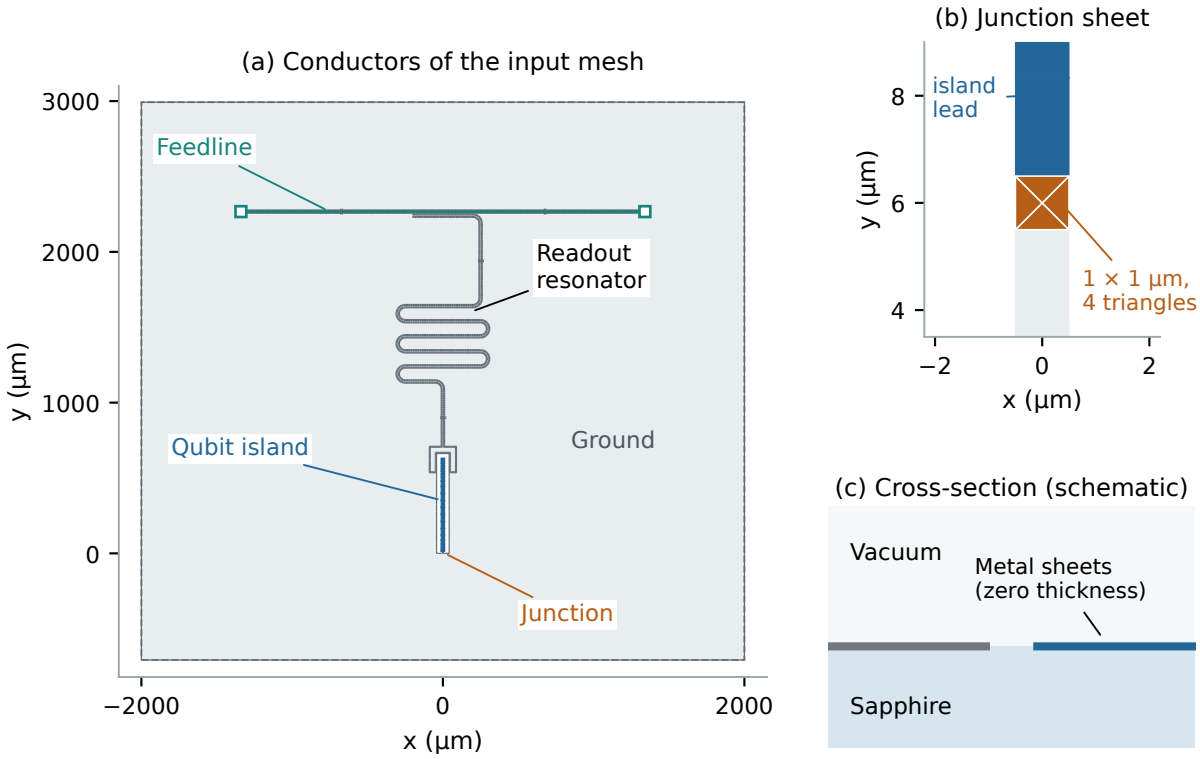


Figure 1: Benchmark geometry, drawn from the committed SHA-256-pinned MSH 4.1 input mesh of the DeviceLayout.jl v1.15.0 `SingleTransmon` device. (a) Top view of the three node-disjoint metal conductors of Section 7: the feedline (teal), the qubit island (blue), and the ground plane together with the meander readout resonator (gray), whose etched gaps are drawn as dark outlines. Squares mark the two readout ports; the dashed line is the exterior-boundary outline. (b) The $1 \times 1\,\mu$m junction `lumped_element` sheet with its actual four-triangle triangulation. (c) Material cross-section, schematic and not to scale; the `substrate` and `vacuum` volume groups fill the enclosing box. Both solvers read this file (22,684 nodes; 133,314 tetrahedra); Palace then refines it internally around the metal sheets, so the two do not solve the same discrete problem (Section 4.1).

self-consistent-but-wrong agreement. Both test the port term's L and C scalings. The first is the resonance of the port sheet's own inductance against its own capacitance, $f_{LC} = 1/(2\pi\sqrt{LC}) = 17.60\,\text{GHz}$, which anchors the 17.49 GHz mode. The second is a $\sqrt{L}$-scaling test: doubling L must move that mode by $1/\sqrt{2}$ — measured: 17.4901 → 12.37 GHz, committed ratio 0.7071 (computed from the unrounded values; Figure 3). Earlier versions called the 17.49 GHz mode "the junction mode" and treated these two checks as pinning the junction physics. Section 6.3 shows that this mode carries no port voltage, so the checks pin the port *operator*, not the qubit; the qubit mode has its own scaling check in Section 6.1. The suite also contains inverse tests that *must fail* (perturbed configurations required to disagree with the oracle — a pipeline that still passes under them is not measuring anything), and the Palace oracle is deterministic: a rerun from the committed configuration reproduced its committed `eig.csv` bit-for-bit.

Table 1: Eigenmode agreement from the identical input mesh (22,684 nodes / 133,314 tets; geode-fem solves 156,863 Nédélec DOFs, 133,108 interior after PEC; Palace refines internally and solves 183,891), first order in both solvers, junction as a lumped reactive shunt ($L = 14.860\,\text{nH}$, $C = 5.5\,\text{fF}$). These are the six modes Palace returns when it searches above 4.5 GHz; the qubit mode below that target is in Table 2. geode-fem at commit `3174015`; Palace at commit `fba6a5b`. All values are quoted verbatim from the committed `benchmarks/transmon_eigen/results.toml`: geode-fem frequencies at the committed three decimals, Palace at full precision, and the committed per-mode `rel_err_pct`, computed against the rounded geode-fem values (precision disclosure in Section 12).

Mode		geode-fem (GHz)	Palace (GHz)	Δ (%)
`resonator`	readout resonator	5.153	5.151335830348	0.032
`mode_2`	box/cavity	15.465	15.46052107794	0.029
`junction_lc`	port-sheet LC*	17.490	17.49010903536	0.001
`mode_4`	box/cavity	18.693	18.69165792915	0.007
`mode_5`	box/cavity	20.703	20.69755679425	0.026
`mode_6`	box/cavity	26.088	26.08089940472	0.027

*The artifact key `junction_lc` is kept verbatim. Palace's port EPR for this mode is 2.49×10^{-8}, and Section 6.3 identifies it as an artifact of the distributed port term.

Mode identification. Modes are identified, not eyeballed. Within geode-fem, the stiffness-participation ratio

$$p \;=\; \frac{x^{\mathsf{T}} K_{\text{port}}\, x}{x^{\mathsf{T}} \left(K + K_{\text{port}}\right) x} \tag{3}$$

measures the share of a mode's stiffness energy stored in the port surface term, whatever its spatial pattern. It is $p = 1.000$ for the 17.49 GHz mode, $p = 0.9942$ for the 3.45 GHz mode, and $p = 0.000$ (at the artifact's three-decimal precision) for the resonator and box modes. Palace's port EPR is built from the port *voltage* functional alone, so it counts only the uniform-voltage part of that energy. The two diagnostics agree where the port field is a uniform voltage and disagree where it is not: on the qubit mode Palace reports 0.99424 against geode-fem's 0.9942, while on the 17.49 GHz mode Palace reports 2.49×10^{-8} against geode-fem's 1.000 (Section 6). In the original Palace run, which searched only above 4.5 GHz, all six port-EPR values are $\leq 5 \times 10^{-4}$; earlier versions read that as two diagnostics ranking the modes differently, when the cause was the qubit mode's absence from Palace's search window. The six-mode comparison of Table 1 is matched by frequency.

Result. Table 1 and Figure 2: all six modes agree to $\leq 0.033\%$ (worst case 0.032%). The 17.49 GHz port-sheet mode agrees most closely, at 0.001%, and sits 0.6% below $f_{LC} = 17.60$ GHz in *both* solvers. Agreement that tight on a mode confined to the four port triangles is what an identical port operator would produce, and it is weaker evidence about the bulk discretizations than the cavity modes' $\sim$0.03%. The qubit mode, which Palace returns only when asked to search below it, agrees to 5×10^{-5} % (Section 6.1, Table 2). The agreement is between genuinely independent code paths — a batched-tensor assembly feeding a direct-factorization Lanczos in Rust, and a projected, preconditioned Krylov method in C++ — on discrete problems that differ by Palace's internal refinement (183,891 against 156,863 unknowns). A common implementation error is therefore an unlikely explanation for the residual. A shared *modeling* property is not excluded by cross-comparison, and Section 6.3 reports one. All of this is same-input-mesh agreement at first order; none of these frequencies is a mesh-converged value.

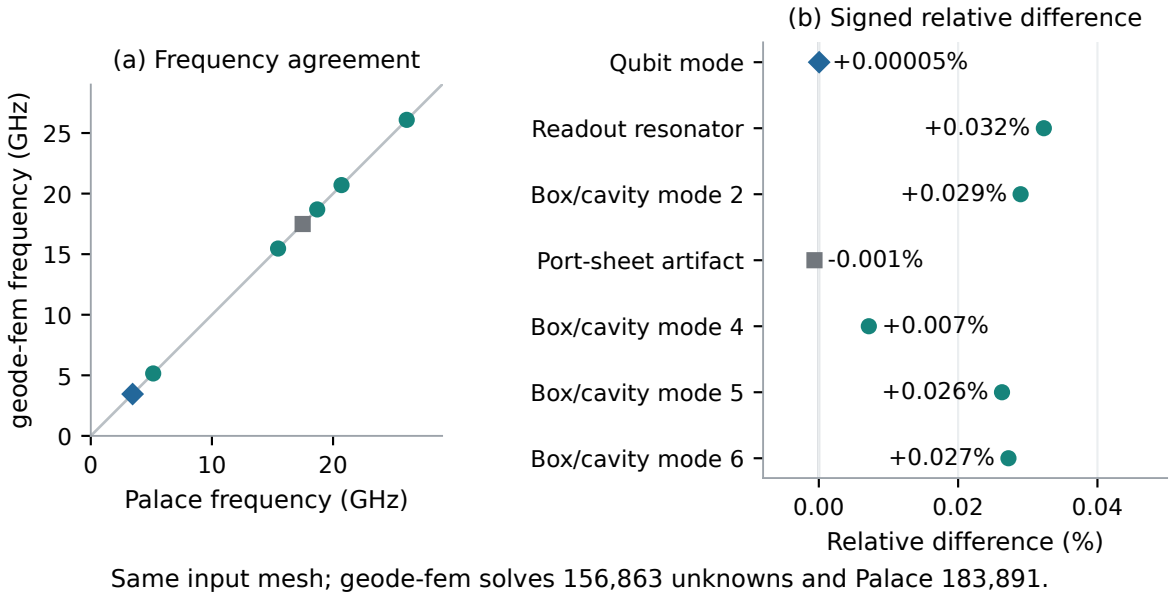


Figure 2: Mode-frequency agreement from the identical input mesh: the six modes Palace returns above its 4.5 GHz search target (Table 1) and the qubit mode from the `Target` = 2 GHz re-run (Table 2). (a) geode-fem against Palace frequency, with the identity line. (b) Signed relative difference $100\,(f_{\text{geode}}/f_{\text{Palace}} - 1)$, computed from the committed frequencies: geode-fem at its committed three decimals for the six modes, which reproduces the per-mode values of Table 1, and at its committed six decimals for the qubit. The 17.49 GHz port-sheet artifact is marked separately; agreement on it tests the shared port operator, not the qubit physics (Section 6.3). Data from `benchmarks/transmon_eigen/results.toml`.

6 Honest Physics: The Qubit Mode, Two Withdrawn Claims, and a Port-Sheet Artifact

Versions 1–6 of this paper made two claims in this section, and both were wrong. The first was that no qubit mode exists in this junction model by construction, and that Palace confirmed its absence. The second was that geode-fem's 3.4528 GHz mode is a spurious port artifact with no Palace counterpart. We withdraw both. The committed evidence identifies the 3.45 GHz mode as the transmon qubit mode, which Palace also returns, and it points to a different mode, at 17.49 GHz, as the artifact of the port model. This section separates what is proven by committed artifacts from what is derived, what a scratch probe measured, and what is only planned.

6.1 The 3.45 GHz mode is the transmon qubit mode

Why the earlier Palace run could not return it. The original Palace run asked for six modes with `Target` = 4.5 GHz and returned nothing below the 5.15 GHz resonator. Earlier versions read that as Palace hunting near 4.5 GHz and finding no qubit mode. That misreads `Target`. Palace's eigensolver accepts only eigenvalues at or above the target: its SLEPc configuration restricts the search region to that side, and its configuration schema defines `Target` as the frequency "above which to search" (`docs/research/transmon-lumped-port-formulation.md`, §2, read from the Palace source at the benchmarked changeset). The run's silence below 4.5 GHz was therefore not evidence of absence.

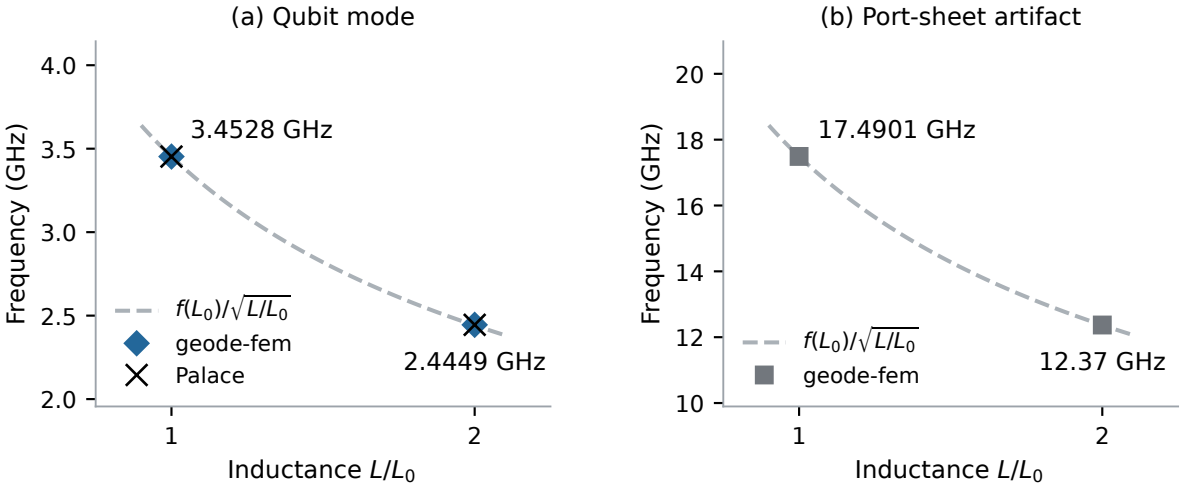


Figure 3: Inductance-doubling checks with $L_0 = 14.860\,\text{nH}$ for (a) the qubit mode and (b) the 17.49 GHz port-sheet artifact. Dashed curves show $f(L_0)/\sqrt{L/L_0}$, anchored at each mode's own baseline. Doubling L moves the port-sheet mode from 17.4901 to 12.37 GHz (committed ratio $0.7071 = 1/\sqrt{2}$), and moves the qubit mode from 3.4528 to 2.4449 GHz in geode-fem (ratio 0.7081) and from 3.452749 to 2.444899 GHz in Palace (ratio 0.70810); at this scale the Palace markers cover geode-fem's. The qubit mode's ratio sits slightly above $1/\sqrt{2}$ because the full-wave mode lies 0.28% below the lumped estimate at L and 0.14% below it at $2L$. Both modes respond to L, so this check alone does not identify the physical qubit mode. Values from the committed `benchmarks/transmon_eigen/results.toml` (`[tripwires.junction_l_doubling]`, `[spurious_mode]` and `[oracles.palace_target2]`).

Proven: Palace returns the mode when asked. We re-ran Palace at the same changeset on the same input mesh, with the same materials, `LumpedPort`, order, eigensolver tolerances and linear solver, changing only `Target` (4.5 → 2.0 GHz) and the number of modes (6 → 8); a second cell also doubles L. A control cell with the original request reproduces the committed spectrum to a relative 6.5×10^{-13}, so the local build does not move the eigenvalues. Table 2 gives the result (`benchmarks/transmon_eigen/results.toml`, `[oracles.palace_target2]`; raw outputs under `reference/fixtures/transmon_palace/results_p1_target2/`).

Palace finds the mode at 3.452749 GHz with port EPR 0.99424; geode-fem's committed value is 3.452751 GHz with $p = 0.9942$. The two frequencies agree to 5×10^{-5} %. Under L doubling the mode moves to 2.444899 GHz in Palace, a ratio of 0.70810, against geode-fem's committed 0.7081. Palace's value lies 0.280% below the lumped estimate at L and 0.140% below it at $2L$. That estimate takes C_Σ from the electrostatic branch on the same coarse mesh, so it is a consistency check between this paper's eigenmode and LOM branches, not an independent anchor. Every other mode in the re-run matches the committed Palace spectrum, and two further modes converge above it (33.384 and 33.896 GHz).

Earlier versions explained the supposed absence by saying that the lumped-shunt model gives the junction only its own 5.5 fF, so that its resonance lands at 17.5 GHz, and that the pad capacitance of a physical transmon (Koch et al., 2007) was missing. That explanation was wrong. The full-wave eigenproblem contains the island's capacitance through the field itself, which is why a mode exists in which the junction inductance resonates against the whole island.

Table 2: The qubit mode in both solvers, from the identical input mesh at first order. Palace values are from the `Target` = 2 GHz re-run; geode-fem values are the committed ones that earlier versions reported as a spurious mode. The lumped estimate is $1/(2\pi\sqrt{L(C_\Sigma + C_J)})$ with the electrostatic $C_\Sigma = 136.684731$ fF of Section 7 and $C_J = 5.5$ fF. Palace's column is its port EPR; geode-fem's is the stiffness participation of Eq. 3. This is a same-input-mesh comparison on a coarse mesh, not a converged qubit frequency.

Quantity	Palace	geode-fem	Lumped estimate
Qubit mode, $L = 14.860$ nH (GHz)	3.452749	3.452751	3.46246
Port EPR / participation p	0.99424	0.9942	—
Qubit mode, $2L$ (GHz)	2.444899	2.4449	2.44833
L-doubling ratio ($1/\sqrt{2} = 0.70711$)	0.70810	0.7081	0.70711
17.49 GHz mode: port EPR / p	2.49×10^{-8}	1.000	—
17.49 GHz mode at $2L$ (GHz)	12.3674	12.37	—

What this does not show. This is not the converged physical qubit frequency. Both solvers start from the same coarse 133k-tetrahedron mesh at first order, and they do not solve the same discrete problem: Palace refines around the metal sheets and solves 183,891 unknowns against geode-fem's 156,863 (Section 4.1). The mesh is also far from converged for the capacitance that sets this frequency: one uniform refinement lowers the scalar-ε C_Σ by 26.3% (Section 7). The lumped estimate goes as $1/\sqrt{C_\Sigma + C_J}$, so we infer that the qubit frequency is not converged on this mesh either; that inference has not been tested by a refined eigensolve. The DeviceLayout.jl blog workflow (Peairs and Carson, 2025) reports a spectrum from 4.14 to 5.591 GHz that includes a ∼4 GHz qubit mode. We do not claim to reproduce it, and we do not attribute the difference from our 3.45 GHz to any single cause.

The committed artifacts still carry the old labels. `benchmarks/transmon_eigen/results.toml` keeps this mode under the section name `[spurious_mode]` with `excluded = true`, and the CPU timing artifacts label it `spurious_port`. The re-baselining of those classifications is tracked separately (issue #1045) and has not been done. We quote the artifact keys verbatim and state here what the evidence now shows.

6.2 The gauge and projection arc, re-read

Earlier versions told the following arc as the diagnosis of a spurious mode. Its measurements stand, every step committed (Section 12); what they mean changes once the two modes are correctly identified. geode-fem's headline Lanczos operates on the raw pencil of Eq. 2 without a gauge projection. (1) Classical tree–cotree elimination (Albanese and Rubinacci, 1988; Manges and Cendes, 1995) is not spectrum-preserving for the generalized eigenproblem — the resonator shifts 1.64% (5.1528 → 5.2372 GHz), outside our bar, pinned as a regression test. (2) The spectrum-preserving bulk divergence-free projector (structurally Palace's confinement) preserves the cavity spectrum (0.029%) and collapses the 13,747-mode gradient cluster. It removes the 17.49 GHz mode, which is almost entirely a gradient field (measured divergence ratio 50.2; projected norm 1.06×10^{-4}), and it keeps the 3.45 GHz mode, which is solenoidal with respect to the interior gradients (divergence ratio $\approx 6 \times 10^{-15}$). Earlier versions described this as the projector deflating the physical junction mode while the spurious mode survived. With the identities corrected it reads the other way: the projector kept the qubit mode and removed the port-sheet mode. (3) A port-aware rank-1 re-admission of one gradient direction restores the 17.49 GHz mode and returns all six of Palace's `Target` = 4.5 GHz modes at $\leq 0.029\%$ worst case. That construction did what it was built to do. Its purpose, though,

was to retain a mode that Section 6.3 identifies as an artifact, because the Palace reference then in hand contained that mode and did not contain the qubit mode.

The remaining committed characterization of the 3.45 GHz mode is unchanged and now reads as a description of the qubit mode: under the L-doubling harness it moves $3.4528 \to 2.4449$ GHz (ratio 0.7081) with 99.4% of its stiffness energy in the K_{port} term. Earlier versions added that Palace's `LumpedPort` represents the port differently and therefore never sees this mode. That was also wrong: Palace's port operator has the same structure (Section 4.2), and Palace sees the mode.

The headline comparison deliberately remains the committed ungauged benchmark. The subsequent port-subspace implementation removes the near-zero survivor and admits the port-patch gradient direction with one shifted factorization. On this fixture, requesting seven modes at 4.5 GHz returns the six modes of Table 1 (worst relative frequency error 0.0289%) plus the 3.452751 GHz qubit mode; requesting six still misses the highest of the six. The committed local comparison records mean process CPU times of 46.32 s for the older eight-mode extraction route and 23.25 s for the new seven-mode route. These are geode-only measurements on a loaded Apple M3 Ultra, not a Palace comparison or a same-box update to the EC2 timings (Section 11.1). geode-fem's benchmark harness still classifies modes by frequency matching against the Palace reference; no reference-free rule is built. Figure 4 shows the spectrum.

6.3 The 17.49 GHz mode: a zero-voltage resonance of the distributed port term

Proven by committed data. Palace's port EPR for the 17.49 GHz mode is 2.49×10^{-8}, in the original run and in the re-run. After L doubling the mode moves to 12.3674 GHz (ratio 0.70711) and its port EPR is 2.94×10^{-8}. Palace builds that EPR from the port voltage, so in Palace this mode carries essentially no port voltage. In geode-fem the same mode has $p = 1.000$: all of its stiffness energy is in the port surface term. Both statements hold only if that energy sits in a surface pattern whose net voltage is zero.

Measured by a scratch probe. A probe on the committed fixture measured the uniform-voltage share of each mode's port energy in geode-fem: 0.00% for the 17.49 GHz mode and 100.00% for the 3.45 GHz mode. The probe source is committed (`docs/research/transmon-lumped-port-formulation.probe-1003.rs`) but it is scratch evidence: the build does not compile it, and no regression test pins these two numbers.

Derived. The junction patch is four triangles with five nodes, of which only the centre node is not on a conductor. The gradient of that node's hat function is curl-free and has zero port voltage, but the isotropic surface term of Eq. 1 still stores energy in it. A field confined to the patch then has Rayleigh quotient $1/(LC)$, which is $f_{LC} = 17.60$ GHz, and the volume mass of the adjacent tetrahedra lowers it to 17.49 GHz. A two-terminal junction cannot carry this pattern. On this reading the mode is a sheet-internal resonance created by representing a lumped element as a distributed isotropic sheet, and both solvers have it because both use that representation. The reading predicts one such mode per free patch node, so a finer port mesh should show a cluster just below f_{LC}. That falsifier has not been run.

What it means for the credential. The 0.001% agreement on this mode is agreement on a shared artifact of the port model. It still tests that the two implementations assemble the same port operator, and nothing more. The physical content of the credential is the five cavity modes and the qubit mode.

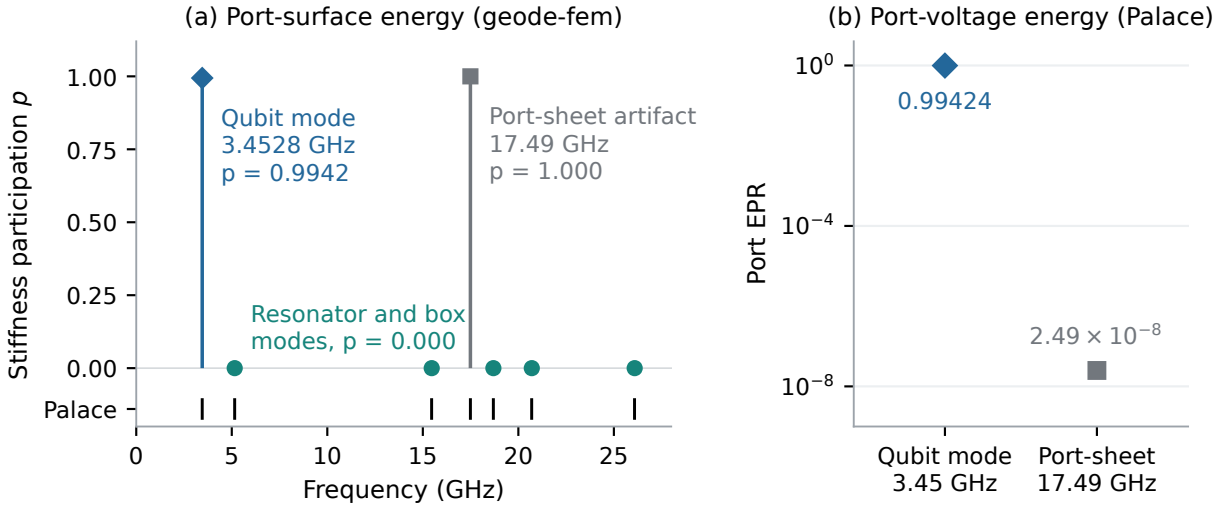


Figure 4: Two different port diagnostics, on separate axes. (a) geode-fem's stiffness participation p (Eq. 3), which counts all port-surface energy, with the mode frequencies of Palace's `Target` = 2 GHz re-run marked below the axis (its two modes above 33 GHz are outside the plotted range). The qubit mode appears at 3.45 GHz in both solvers. It and the 17.49 GHz port-sheet artifact both have participation near one in geode-fem ($p = 0.9942$ and $p = 1.000$); the resonator and box modes have $p = 0.000$. (b) Palace's port EPR, built from the port voltage, on a logarithmic axis: 0.99424 for the qubit mode but 2.49×10^{-8} for the artifact, whose port energy carries no port voltage (Section 6.3). The two quantities coincide only for uniform-voltage port fields (Section 5), so they are not interchangeable measures of participation.

Roadmap, not built. The planned remedy is a rank-1 port-voltage formulation (issue #1044): replace S_Γ in Eq. 1 by $b\,b^\mathsf{T}$, with $b_i = (1/w)\int_\Gamma N_i \cdot \hat{e}\, dS$ the port-voltage functional, so that only the energy of the port voltage is kept. In the same scratch probe this formulation removes the 17.49 GHz mode (no mode remains between 15.47 and 18.69 GHz), leaves the 3.4528 GHz mode and the five cavity modes unchanged to four decimals, and makes the plain bulk projector exact (six modes from a six-mode request, residuals $\leq 1.7 \times 10^{-9}$, no reference filter). These are probe measurements on uncommitted solver code. No shipped solver path, committed benchmark or table in this paper uses the formulation. A solver using it would no longer reproduce Palace's 17.49 GHz mode, and whether it becomes the default is an open decision (issue #1046).

7 The Forward LOM Chain: From Fields to E_C

The lumped-oscillator-model (LOM) branch computes the transmon's Hamiltonian parameters from the *electrostatic* problem on the same committed mesh — and it is the chain the remainder of this paper differentiates. The mesh's metal is split into three node-disjoint conductors (ground+resonator, feedline, island); a tensor-ε electrostatic solve yields the Maxwell capacitance matrix by the energy-reaction form $C_{ij} = \phi_i^\mathsf{T} K \phi_j$; the island self-capacitance is reduced for the floating feedline ($C_\Sigma = C_{ii} - C_{if}^2/C_{ff}$); and the charging energy follows as $E_C/h = e^2/(2C_\Sigma h)$, with the transmon spectrum evaluated charge-basis-exactly after Koch et al. (2007).

Measured on the committed fixture (`benchmarks/transmon_quantum/results.toml`): $C_\Sigma = 136.7\,\text{fF}$, hence $E_C = 0.142\,\text{GHz}$, $E_J/E_C = 77.6$ (with E_J from the junction inductance), $\omega_{01} = 3.38\,\text{GHz}$, and anharmonicity $\alpha = -0.158\,\text{GHz}$, with the $\alpha \approx -E_C$ sanity holding. The chain

Table 3: Mesh convergence of C_Σ (fF) for the unmodified device under uniform refinement of the committed mesh. Every entry is a conforming energy minimum and therefore an upper bound on the converged value for the committed polygonal geometry. No P2 solve was run at $h/4$, and the tensor-ε model has no P2 path. Source: `benchmarks/transmon_diffopt/c_sigma_convergence.toml`.

Level	Tetrahedra	P1, scalar ε	P2, scalar ε	P1, tensor ε
h	133,314	137.706752	90.044141	136.684731
$h/2$	1,066,512	101.359397	83.015285	102.375301
$h/4$	8,532,096	87.753422	n/a	89.493342

runs end-to-end from committed geometry. These are coarse-mesh numbers and they are not mesh-converged, as the next paragraph shows. The linear eigenproblem resolves the same oscillator at 3.45 GHz (Section 6.1); ω_{01} is its anharmonic $0 \to 1$ transition, computed from a C_Σ that omits the junction's own 5.5 fF, so the two numbers are not expected to coincide.

Mesh convergence of C_Σ. A convergence study on the committed geometry (`benchmarks/transmon_diffopt/c_sigma_convergence.toml`) refines the committed mesh uniformly, splitting every tetrahedron into eight, and re-solves C_Σ of the unmodified device, written X^0 below. It uses the scalar-ε model, the one the shape gradient differentiates, with linear (P1) elements at h, $h/2$ and $h/4$ and quadratic (P2) elements at h and $h/2$ (Table 3, Figure 5). The $h/4$ mesh has 8,532,096 tetrahedra. A sparse direct factorization does not fit safely at that size on the 96 GiB host, so $h/4$ and P2 at $h/2$ are solved by conjugate gradients with an algebraic-multigrid preconditioner; where both solvers run, they agree to a relative 10^{-9} or better.

Rigorous. The P1 spaces are nested under the same Dirichlet and floating constraints, the P2 space on a mesh contains the P1 space on it, and C_Σ is an energy minimum. Every solved value is therefore an upper bound on the converged C_Σ of this geometry. The smallest, P2 at $h/2$, puts the converged scalar-ε value of X^0 below 83.02 fF. The coarse 137.706752 fF overstates it by at least 54.69 fF, which is at least 65.8% of the converged value. There is no rigorous lower bound.

Estimate, not a bound. Four extrapolations span 74.14–79.62 fF: Richardson extrapolation of the P1 levels at the observed order and at order one, a fit of $C_\infty + ah + bh^2$ through the three P1 levels, and Richardson extrapolation of the two P2 levels at order one. The observed P1 order is 1.41. We expect order one, because the potential near the edge of a zero-thickness sheet behaves like $r^{1/2}$; that expectation is an inference from theory and not a measurement. The three levels are therefore pre-asymptotic. The band is not asymptotic and it is not a bound. The P2 value on the committed mesh, 90.044141 fF, happens to fall near the anchor. It is not converged either: one refinement lowers it by a further 7.02 fF.

Tensor permittivity. With the sapphire tensor, the P1 values in Table 3 fall in the same way. The converged tensor C_Σ is below 89.50 fF, and two Richardson values put it at 76.61–81.75 fF, again an estimate. The scalar-minus-tensor difference changes sign under refinement: it is +0.75%, −0.99% and −1.94% of the tensor value at h, $h/2$ and $h/4$. The 0.75% scalar-versus-tensor difference is thus a coarse-mesh number and not a converged model difference.

Scope. "Converged" means converged for the polygonal surfaces of the committed mesh, which refinement does not move. It is not a statement about a regenerated or re-meshed DeviceLayout device.

Consequences for the chain. Every C_Σ, E_C, E_J/E_C, ω_{01} and α quoted at the start of this section is a coarse-mesh number. The coarse tensor C_Σ overstates the converged one by at least 47.18 fF, and because $E_C \propto 1/C_\Sigma$ the converged E_C is at least 52.7% above the coarse 0.142 GHz. The Koch

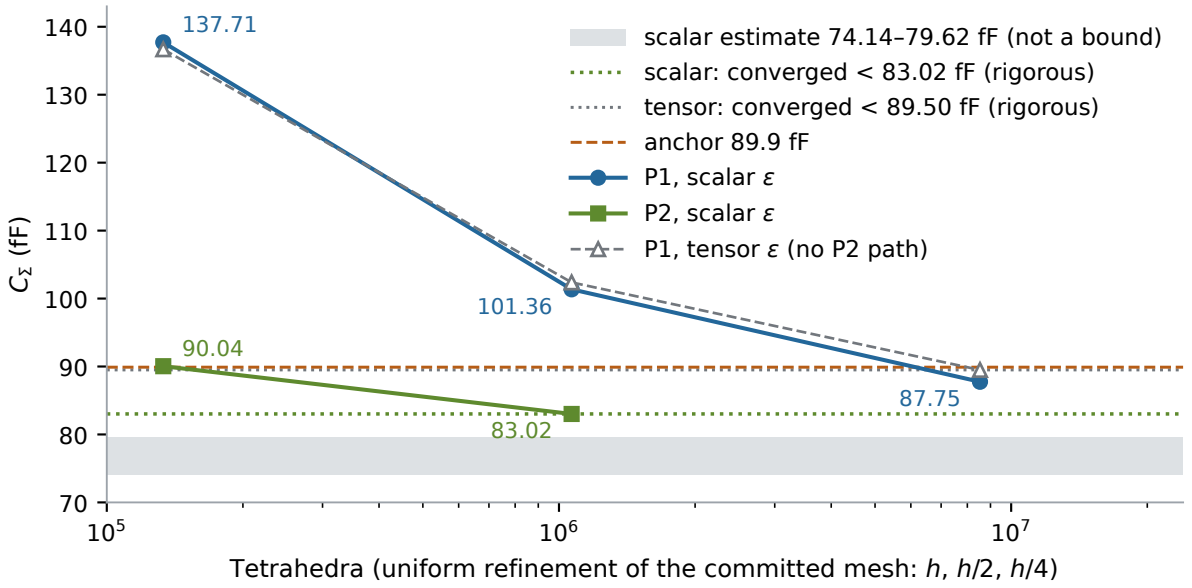


Figure 5: Mesh convergence of C_Σ for the unmodified device X^0 under uniform refinement of the committed mesh (Table 3). Every plotted point is a conforming upper bound on the converged value of the committed polygonal geometry; the shaded band is an extrapolated estimate and is not a bound. Scalar permittivity: P1 at h, $h/2$ and $h/4$ and P2 at h and $h/2$, with the rigorous upper bound of 83.02 fF and the estimate band of 74.14–79.62 fF. Tensor permittivity, which has no P2 path: P1 only, converged below 89.50 fF. The dashed line is the 89.9 fF anchor. The held-out refinement of the two optimized designs is Figure 6b. Source: `benchmarks/transmon_diffopt/c_sigma_convergence.toml`.

spectrum has not been re-evaluated at a converged C_Σ.

The gap to the design anchor was discretization error. The coarse $C_\Sigma = 136.7$ fF *exceeds* the ∼90 fF design-anchor value back-solved from the blog workflow's expected $E_C = 0.2156$ GHz. Earlier versions, following the benchmark as then committed, diagnosed this as a model-definition difference and not a truncation artifact, on two checks: grounding the exterior wall moves C_Σ by a relative 6×10^{-5}, and the scalar-versus-tensor sapphire treatment shifts it 0.75% on the coarse mesh. Neither check tests mesh convergence, and the diagnosis is withdrawn. The convergence study settles the question. The converged scalar-ε C_Σ of X^0 is below 83.02 fF, which is at least 6.88 fF, or 7.6%, *below* the 89.9 fF anchor; the estimate band puts it 11.4–17.6% below. The coarse mesh overstates C_Σ by at least 54.69 fF, more than the whole 47.8 fF by which the coarse value exceeds the anchor. The gap was discretization error, and more than all of it.

What remains is a gap of the opposite sign: converged, geode-fem's electrostatic C_Σ is below the anchor. As an inference only, the junction's 5.5 fF self-capacitance, which is outside this C_Σ, closes part of it: added to the estimate band it gives 79.64–85.12 fF, and added to the rigorous bound it gives 88.52 fF, both still below the anchor. The rest is not explained here; candidates are the geometry representation and the anchor's own derivation. We report the measured values and the discrepancy; we do not retrofit to the anchor. Section 9 uses the anchor as an optimization target and reports what that does and does not establish.

Table 4: The merged 2×2 discrete-adjoint sensitivity matrix. Each cell: one forward + one adjoint solve sharing a single factorization; element Jacobians exact (forward-mode duals through the production kernels); validated against central finite differences of the full pipeline, with mutation tests. "FD rel-err" is the observed agreement of the analytic gradient with the FD reference in the merged test suite (test acceptance bounds are looser; see the module sources). Modules in `crates/geode-core/src/`.

Operator	Parameter	Module (issue / PR)	FD rel-err
scalar electrostatic (SPD)	material $\partial/\partial\varepsilon$	`adjoint` (#570 / #573)	$\sim 3 \times 10^{-8}$
scalar electrostatic (P1)	geometry $\partial/\partial X$	`shape` (#571 / #575)	$\sim 1 \times 10^{-9}$
driven $H(\mathrm{curl})$	material $\partial/\partial\varepsilon$	`driven::adjoint` (#576 / #579)	2.3×10^{-5} (worst region)
driven $H(\mathrm{curl})$	geometry $\partial/\partial X$	`driven::shape` (#577 / #581)	2.2×10^{-9}, 1.2×10^{-8} (two maps)

8 Differentiable Sensitivities: The Discrete-Adjoint Layer

8.1 The 2×2 sensitivity matrix

For a solve $A(\varepsilon, X)\, x = b(X)$ and a scalar observable $g(x)$, the discrete adjoint $A^{\mathsf{T}}\lambda = \partial g/\partial x$ turns the parameter gradient into explicit contractions:

$$\frac{dg}{d\varepsilon_k} = -\lambda^{\mathsf{T}} \frac{\partial A}{\partial \varepsilon_k} x, \qquad \frac{dg}{dX} = \lambda^{\mathsf{T}} \frac{\partial b}{\partial X} - \lambda^{\mathsf{T}} \frac{\partial A}{\partial X} x \tag{4}$$

(real case; the complex driven case takes $2\,\mathrm{Re}[\cdot]$ of the same contractions with the *transpose*, not conjugate-transpose, solve — the pencil is complex-symmetric, so the adjoint solve reuses the forward LU factorization). Everything on the right-hand side is assembly-level and therefore differentiable on the tensor substrate: the element Jacobians $\partial A/\partial X$ and $\partial b/\partial X$ are *exact*, computed by forward-mode dual numbers through the same closed-form element kernels the solver assembles — no finite-difference truncation enters the gradient. Each gradient costs one forward plus one adjoint solve sharing a single factorization (asserted programmatically: `n_factorizations` $= 1$).

Four adjoint modules are merged and validated, spanning material and geometry sensitivities on both operators (Table 4). Each carries a committed test that the analytic gradient match a central finite difference *of the full independent pipeline* (perturb the parameter, re-assemble, re-solve through the public solver path, recompute g) — and, load-bearingly, mutation tests proving the tolerances bite: the classic complex-adjoint conjugation error (A^{H} for A^{T}) is rejected by the FD check, as is a dropped term. The driven-shape case surfaced a physics subtlety worth recording: the current-source right-hand side $b = \mathrm{i}\omega\mu_0 \int N \cdot J$ is geometry-dependent, so the shape gradient carries a $\partial b/\partial X$ term absent from the material case — omitting it fails the finite-difference cross-check outright.

The matrix is deliberately 2×2: material and geometry are the two design-parameter classes, and the scalar and driven operators are the electrostatic (LOM) and RF (S-parameter/EPR-workhorse) branches of the solver. The coarse-mesh LOM tiers of Section 9 exercise the scalar column end-to-end, and the refined optimization of Section 9.4 uses the same stationarity identity in a native-f64 P2 implementation; the driven column is merged and validated at the operator level and composes with the driven solve, the one solve geode-fem's GPU path runs (correctly, but slower than both assembled CPU solvers at every measured size, and at best at parity with the same matrix-free algorithm on the CPU; Section 11.3) — RF shape optimization is the natural follow-on, not exercised here.

8.2 The capacitance$\to E_C$ chain, and a vanishing adjoint

Composing the shape adjoint with the LOM chain of Section 7 gives $\partial(C, E_C, \alpha)/\partial(\text{geometry})$ (merged as `shape::capacitance_shape_gradient`, #583 / PR #586). The composition carries an

elegant simplification. The capacitance observable is the stored field energy, $C = 2W = \phi^{\mathsf{T}} K(X)\phi$, and W is *variationally stationary* in the solved potential ϕ: at the forward solution the free rows of the cotangent $K\phi$ vanish (equilibrium), so the adjoint solution $\lambda \approx 0$ and the entire shape derivative collapses to the explicit-geometry term

$$\frac{dC}{dX} \;=\; \phi^{\mathsf{T}} \frac{\partial K}{\partial X}\,\phi \qquad \text{(at the solution)}, \tag{5}$$

a Hellmann–Feynman-like result: no adjoint back-solve is needed for the capacitance itself. The implementation keeps the general adjoint form (the mutation test measures the implicit part at $< 10^{-8}$ of the total, confirming the collapse numerically rather than assuming it), and chains through $\partial E_C/\partial C_\Sigma = -e^2/(2C_\Sigma^2 h)$ exactly. Validation is double: against the analytic parallel-plate law ($\partial C/\partial d = -\varepsilon_0\varepsilon_r$ per unit area under a gap scaling; $\partial C/\partial A = +2\varepsilon_0\varepsilon_r$ under an area scaling — opposite signs, both matched) and against central finite differences of the full pipeline, with observed relative errors at the 10^{-9} level.

9 Gradient-Based Optimization: A Loop Proof, a Coarse-Mesh Detour, and a Refinement-Validated Design

Four tiers, labeled by what each one establishes: a closed-loop proof on an analytic fixture (Section 9.1); single-parameter runs on the coarse committed device mesh, which validate the gradient and are otherwise historical (Section 9.2); a three-parameter optimizer on that coarse mesh, kept as a cautionary precursor (Section 9.3); and the real-device result, a three-parameter optimization on an adaptively refined mesh with quadratic elements, followed by a held-out refinement series that bounds and estimates what the design is worth when converged (Section 9.4). The first two tiers use damped Newton on $E_C(\theta) - E_C^{\text{target}}$ (equivalently $C(\theta)$) with the analytic $dE_C/d\theta$ from Section 8.2; each Newton step costs one forward plus one adjoint solve on a single factorization.

One distinction governs the device tiers. The convergence study of Section 7 shows that the coarse committed mesh overstates C_Σ by at least 54.69 fF, and that the converged device lies *below* the 89.9 fF anchor, not above it. The coarse-mesh tiers were run before that was known. Their gradient validations and convergence statements are exact statements about the coarse discrete problem: the finite-difference checks test the derivative of the discrete C_Σ, and the fresh solves test self-consistency on one discretization. None of them is evidence about the converged device, and the direction they moved in, lowering C_Σ toward the anchor, is the wrong one for it. Only Section 9.4 measures what happens to a design under refinement. The framing claim is capability and step-count, not wall clock: a derivative-free loop (the parameter sweeps of current practice) spends N_{params} extra forward solves per step to estimate what the adjoint delivers exactly for one. No timing comparison against HFSS/Palace-based optimizers is measured or claimed.

9.1 Loop proof: Newton to a target E_C on a parallel-plate fixture

On a parallel-plate capacitor fixture ($\varepsilon_r = 3$; 343 nodes, 1,296 tets; θ scales the gap, so $d = L(1+\theta)$), Newton iteration drives the design from $C_0 = 120.0$ fF ($E_C = 0.1614$ GHz) to the design-anchor target $E_C = 0.2156$ GHz. Full Newton converges in *one* step — and we state the reason rather than dress it up: $C \propto 1/(1+\theta)$ makes $E_C(\theta)$ *affine* in θ, so one-step convergence is expected, and this tier proves the loop end-to-end (gradient, chain rule, update, convergence test), not a hard optimization. The converged $\theta = 0.335657909$ matches the closed-form optimum to 3×10^{-15}. The honesty check that matters: a *fresh, independent* forward solve at the converged geometry — assemble, solve, extract, no adjoint machinery — reproduces the target E_C to a relative error of 1.4×10^{-15}. A damped run ($\alpha = 0.5$) converges in 13 steps and supplies the descent illustration (Figure 8a,

Appendix A). At the hit target, the Koch-exact spectrum gives $E_J/E_C = 51.0$, $\omega_{01} = 4.128\,\text{GHz}$, $\alpha = -0.247\,\text{GHz}$ — the $\sim$4 GHz qubit regime the design anchor implies. All trajectory values are committed in `benchmarks/transmon_diffopt/results.toml`.

9.2 The coarse device mesh: validated gradients and a shortfall that was a mesh artifact

The second tier runs on the coarse committed 133k-tet SingleTransmon mesh, with a fixed-topology shape parameter θ that scales the island conductor in-plane about its centroid (101 island nodes move; the mesh connectivity is fixed). Two results stand as statements about the coarse discrete problem, and a third is historical (`benchmarks/transmon_diffopt/pad_results.toml`):

The gradient is real on the device mesh. At $\theta = 0$ the adjoint gradient $\partial C_\Sigma/\partial\theta = 198.198\,\text{fF}$ is validated against a central finite difference of the *full independent pipeline* (move the island nodes, re-assemble, re-solve, re-extract): relative error 1.15×10^{-4} at $h = 10^{-4}$, with a clean $O(h^2)$ truncation sweep ($1.2 \times 10^{-2} \to 1.0 \times 10^{-3} \to 1.15 \times 10^{-4} \to 1.0 \times 10^{-5}$ as h falls $10^{-3} \to 3 \times 10^{-5}$). This was the first real-geometry validation: the same machinery that passed at 10^{-9}-level on analytic fixtures holds up on production geometry. It is a statement about the coarse discrete problem, and the refined-mesh checks of Section 9.4 supersede it as the load-bearing one.

The loop genuinely converges on the device mesh. A demonstration target within the mesh's validity budget ($C_\Sigma = 137.0\,\text{fF}$) converges in two Newton steps — genuinely non-affine now (the gradient changes by $\sim$21% along the path) — and the converged design is confirmed by a fresh, independent multi-conductor extraction at relative error 5.6×10^{-6} (Figure 8b). The entire coarse-mesh run (validation sweep plus both optimization runs) takes 10.7 s.

Historical: the anchor "shortfall" was an artifact of an unconverged mesh. On the coarse mesh the unmodified device has $C_\Sigma = 137.7\,\text{fF}$, so the anchor of Section 7 (89.9 fF; $E_C = 0.2155\,\text{GHz}$ — the last digit differs from the 0.2156 GHz quoted elsewhere because this run targets the rounded 89.9 fF rather than the back-solved 89.843 fF; both values are artifact-faithful) looked like a reduction of 47.8 fF. Two single-parameter attempts at it stalled on mesh validity, and earlier versions reported them as the paper's honest negative.

Rigid island scale. The $\theta = 0$ gradient puts the anchor at $\theta \approx -0.241$. The fixed-topology map inverts its first tetrahedron at $\theta = -0.0097$, and the run's conservative budget (minimum tet-volume ratio 0.25) caps $\theta_{\text{safe}} = -0.0073$, where the bounded Newton run stalls at $C_\Sigma = 136.5\,\text{fF}$. By a local linear estimate that is 33× short.

Harmonic mesh motion. Extending the prescribed island velocity into the volume with a P1 Laplace solve spreads the deformation. The field D is computed once (0.34 s; the whole run takes 7.8 s), giving $X(\theta) = X_0 + \theta D$. At the same volume-ratio threshold the safe negative step grows from -0.007258 to -0.064609, an 8.90× extension. The initial derivative is 209.087 fF per unit θ, with a full-pipeline central-FD relative error of 1.054×10^{-4}. Bounded Newton reaches 128.792998 fF, confirmed by a fresh extraction, and stalls at the safe boundary, leaving 38.892998 fF to the anchor. The artifact's initial-slope extrapolation gives $\theta_{\text{anchor}} = -0.2286$, or 3.54× the harmonic budget. The sensitivity falls along the measured path, from 209.1 to 114.5 fF per unit θ, so *if the sensitivity does not recover beyond the measured endpoint* the total is at least 6.25× the safe budget (computed 6.256959, floored). That is a conditional extrapolation, not a measured bound.

These are correct statements about the coarse discrete problem, and the mesh-validity limit of a fixed-topology morph is real. The shortfall was not. It was an artifact of an unconverged

mesh: converged, the unmodified device is already below the anchor (Section 7), so the 47.8 fF reduction these runs attempted is not one the device needs, and the 33× and 6.25× factors describe no physical design problem. The refined optimization of Section 9.4 moves the other way, by about 11.5 fF, and stays well inside the same volume-ratio floor. Both coarse runs use scalar trace-averaged sapphire.

9.3 A cautionary precursor: three parameters on the coarse mesh

Before the convergence study, a three-parameter optimizer was run on the coarse mesh (`benchmarks/transmon_diffopt/multiparam_gradient.toml`, `benchmarks/transmon_diffopt/multiparam_results.toml`). Three harmonic morph fields change the island length, the pad width and the cutout gap; each is exactly zero on the junction nodes and the island leads, so no parameter moves the junction region. Their gradients agree with full-pipeline finite differences to relative errors of 8.0×10^{-7}, 3.4×10^{-10} and 1.2×10^{-8} (rounded up). A bounded Gauss–Newton optimizer with per-step volume-stiffened re-morphing then drove the coarse-mesh scalar-ε C_Σ from 137.706752 down to 89.900001 fF in 18 steps, within a 10 kHz tolerance on E_C and with fresh-solve agreement of 1.39×10^{-14}. The stiffening exponent was selected post hoc, and a neighboring exponent reaches a different geometry in 14 steps, so the design is not unique.

Every one of those checks passed, and as a device design the result is wrong: *the optimizer moved the wrong way.* It lowered C_Σ by 47.8 fF, whereas the converged device starts below the anchor and needs an increase. Under refinement the capacitance of that design falls to 66.609133 fF (P1, $h/2$), 56.298812 fF (P1, $h/4$) and 52.368047 fF (P2, $h/2$) (`benchmarks/transmon_diffopt/c_sigma_convergence.toml`). Its converged value is therefore below 52.37 fF, at least 41.7% below the anchor, with an extrapolated estimate of 45.98–48.11 fF. The lesson is methodological. Finite-difference checks and fresh solves certify consistency on one discretization and say nothing about discretization error. What survives from this tier is its machinery: the junction-pinned morph fields, which the refined optimization reuses.

9.4 The real device: a refinement-validated optimization

This is the paper's real-device design result. Its scope is the one the committed artifacts state (`benchmarks/transmon_diffopt/refined_optimization/`): the scalar-ε electrostatic island capacitance with a neutral floating feedline, on the polygonal geometry of the committed mesh, with the junction's 5.5 fF self-capacitance excluded. It is not an optimization of the full-wave eigenfrequency or of the nonlinear Hamiltonian.

Method. Ten cycles of adaptive refinement, marked by a flux-jump indicator and solved with P1 elements, turn the committed mesh into a 448,279-tetrahedron mesh. The optimizer then works on that mesh with P2 fields, solved by conjugate gradients with an algebraic-multigrid preconditioner. The geometry map is $X(\theta) = X^0 + \sum_i \theta_i D_i$, with the three junction-pinned harmonic fields of Section 9.3 (island length, pad width, cutout gap) computed on the coarse mesh and carried to each refined mesh by exact nodal prolongation. Connectivity, conductor membership and material values stay fixed, and refinement happens outside the differentiable map. The gradient is an equilibrium-energy shape derivative. For a unit island excitation the floating feedline sits at $-C_{fi}/C_{ff}$, the energy of the effective field $\phi_i - (C_{fi}/C_{ff})\,\phi_f$ is C_Σ, and stationarity removes the implicit field derivatives, as in Eq. (5). For a morph velocity v the P2 derivative integrates $\nabla\cdot v\,|\nabla\phi|^2 - 2\,\nabla\phi^{\mathsf{T}}\,\nabla v\,\nabla\phi$ with exact four-point quadrature on each affine tetrahedron. This numerical path is native f64 and does not run on the Burn tape. Accepted steps must reduce the freshly solved capacitance residual and keep the minimum element-volume ratio above 0.25.

Table 5: Held-out refinement of the optimized designs. Round 0 is the optimization mesh; no later round feeds the optimizer. "Design" is the primary run, targeted at 89.9 fF; X^0 is the unmodified device on the same meshes, and the increment is their difference. "Corrected" is the resolution-corrected design on its own meshes. All capacitances are in fF with P2 elements, and each is an upper bound on the converged value of its geometry. Source: the two `summary.toml` files under `benchmarks/transmon_diffopt/refined_optimization/`.

Round	Tetrahedra	Design	X^0	Increment	Tetrahedra	Corrected
0	448,279	89.900518	78.387918	11.512600	448,279	90.980104
1	528,324	89.662755	78.177697	11.485058	528,677	90.736602
2	684,646	89.470366	78.011928	11.458438	684,107	90.538648
3	900,226	89.324261	77.882041	11.442221	905,316	90.391037
4	1,239,090	89.216843	77.786774	11.430069	1,247,811	90.281356
5	1,693,481	89.128359	77.710173	11.418186	1,708,342	90.191606
6	2,369,256	89.065595	77.653890	11.411705	2,398,139	90.127899
7	3,289,543	89.014998	77.609945	11.405054	3,331,618	90.076691

Primary run. On the adapted mesh the unmodified device has $C_\Sigma = 78.387918$ fF, below the anchor. The optimizer therefore *increases* C_Σ, and reaches 89.900518 fF in three accepted steps (Figure 6a; `benchmarks/transmon_diffopt/refined_optimization/p2_cycles10/summary.toml`). The final geometry has island length 627.323864 μm and pad width 29.087261 μm, from 625.5 and 24 μm (Figure 6c–e). The minimum volume ratio over the accepted steps is 0.61413. At the final design all three sensitivities are checked against central differences of two fully re-solved perturbed geometries; the worst relative discrepancy is 3.87×10^{-7}. The target is reached on that discretization, and only there.

Held-out refinement. The final geometry is then refined for seven further rounds with the same marking rule, with a fresh P2 solve each round. None of it feeds the optimizer. Table 5 and Figure 6b give the series. The capacitance falls monotonically from 89.900518 fF at 448,279 tetrahedra to 89.014998 fF at 3,289,543, and it is still falling, by 0.05 fF in the last round.

Rigorous. The held-out spaces are nested and conforming and the capacitance is an energy minimum, so every value is an upper bound on the converged capacitance of this polygonal geometry. The converged value is below 89.02 fF, which is at least 0.98% below the anchor. There is no rigorous lower bound.

Estimate, not a bound. The decrement ratios show no trend toward zero. A geometric tail at the largest and the smallest observed ratio gives 88.778547 to 88.891530 fF, and least-squares fits of $C_\infty + aN^{-\alpha}$ with a free rate on three windows of rounds give 88.900718, 88.863790 and 88.892726 fF. The outward-rounded union is 88.77–88.91 fF, about 88.8–88.9 fF. On that estimate the optimized 89.900518 fF overstates the converged value by 1.11–1.28%. The estimate depends on the assumed rate: the same fit at slower assumed rates gives 88.336686 fF ($\alpha = 1/3$) and 88.618200 fF ($\alpha = 1/2$), below the band. That is why the band is an estimate and not a lower bound.

No 1% claim. The rigorous statement, at least 0.98% below the anchor, does not decide whether the design is within 1% of it. The estimate puts it 1.10–1.26% below. We do not claim the 1% threshold is met, and we infer from the extrapolation that it is not.

The design increment is stable. On the same meshes the unmodified device goes from 78.387918 to 77.609945 fF, so its converged value is below 77.61 fF, at least 13.6% below the anchor, and is estimated at 77.42–77.52 fF. That is inside the 74.14–79.62 fF band of the uniform study and tighter than its 83.02 fF bound (Section 7). The increment the optimizer added moves much less than either capacitance: it is 11.512600 fF on the optimization mesh and 11.405054 fF after seven rounds, a

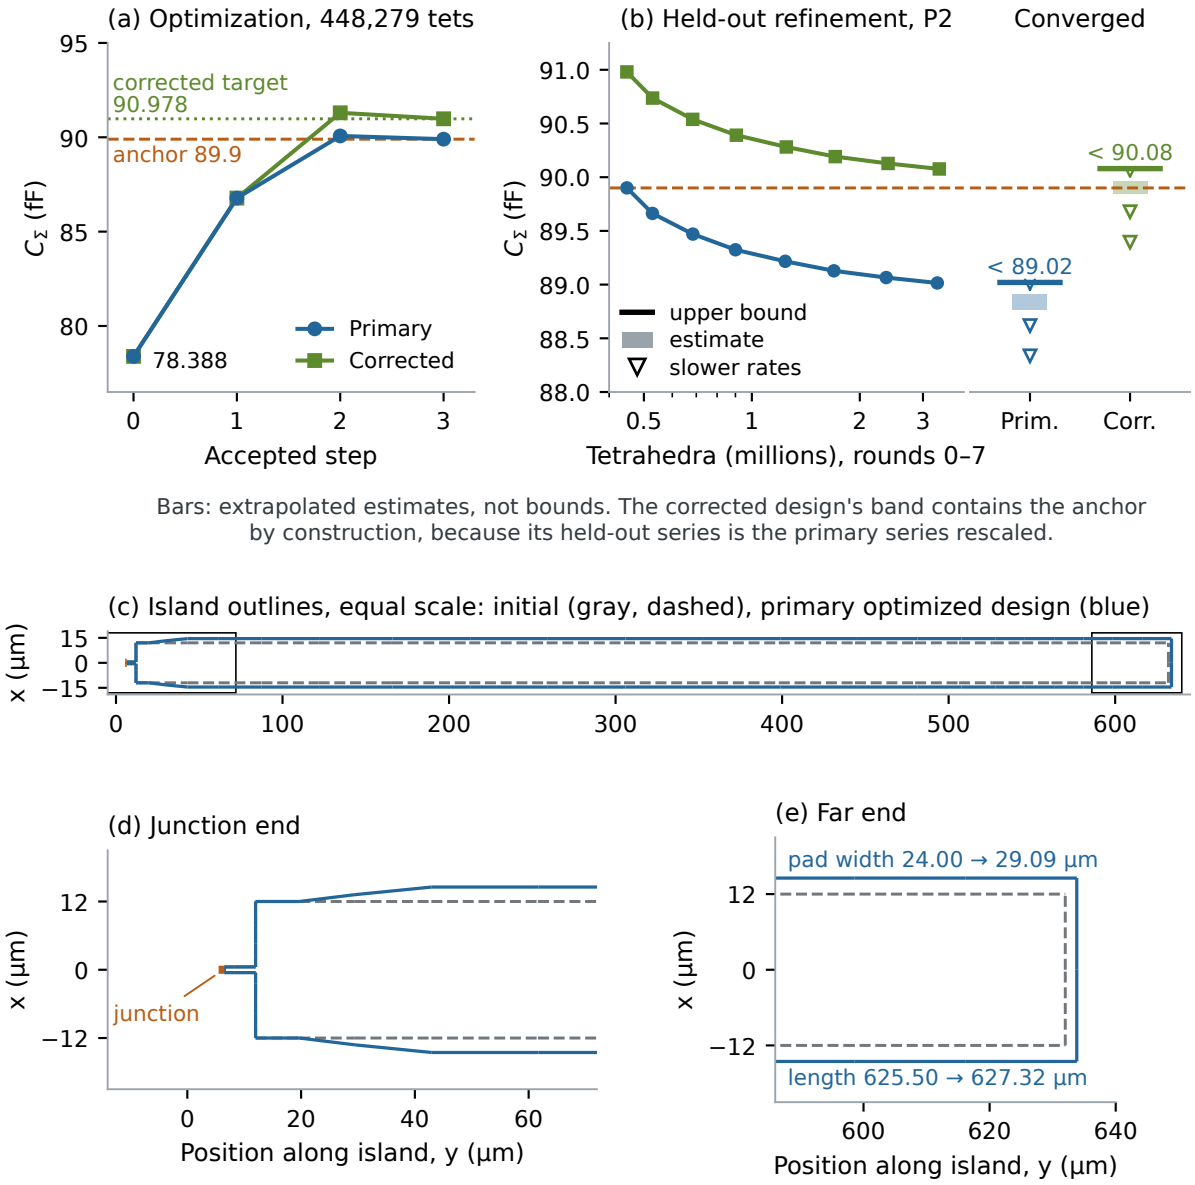


Figure 6: The refined real-device optimization of Section 9.4: the scalar-permittivity electrostatic C_Σ, with a neutral floating feedline and without the junction's 5.5 fF, from P2 solves on the committed polygonal geometry. (a) Freshly solved C_Σ after each accepted step on the 448,279-tetrahedron optimization mesh, for the primary run (target: the 89.9 fF anchor) and the resolution-corrected run (target 90.978253 fF). Both start from the unmodified device, below the anchor, and increase C_Σ. A target reached here is reached on this discretization only. (b) The seven held-out refinement rounds of each final design (Table 5); none feeds the optimizer. Every point is an upper bound on the converged value of its geometry: the primary design converges below 89.02 fF and the resolution-corrected design below 90.08 fF. The bars at the right are extrapolated estimates, not bounds (88.77–88.91 fF and 89.84–89.96 fF), with the same fit at the two slower assumed rates marked beside each. The corrected design's estimate band contains the anchor by construction, because its series is the first one rescaled. No lower bound is proven for either design. (c) Initial and optimized island outlines at equal axis scales, from the committed input mesh and from the same mesh moved to the primary design as exported by that run (its polygonal surfaces are those of the adapted mesh); (d) and (e) enlarge the two ends. The junction lead does not move. Source: the two `summary.toml` files under `benchmarks/transmon_diffopt/refined_optimization/`.

change of less than 1%. Most of the discretization error is common to the two geometries on these meshes, which are adapted to the final design and not to X^0. That is an observation about this series, not a bound.

A resolution-corrected design, and what it does not show. A second run repeats the optimization on the same discretization with the target offset by the measured bias: the anchor times the optimization-mesh value over the geometric-tail midpoint, $89.9 \times 89.900518/88.835038 = 90.978253\,\mathrm{fF}$. The first run's artifact records that target, which was fixed before the second run (`benchmarks/transmon_diffopt/refined_optimization/p2_cycles10_resolution_corrected/summary.toml`). The optimizer reaches 90.980104 fF in three accepted steps (island length 627.452356 μm, pad width 29.447205 μm, minimum volume ratio 0.57790, worst finite-difference discrepancy 3.94×10^{-7}). Its held-out series runs from 90.980104 to 90.076691 fF at 3,331,618 tetrahedra.

Rigorous. The converged capacitance of this geometry is below 90.08 fF, at most 0.20% above the anchor. The last solved value is still above the anchor, so nothing rigorous places the converged value near it from below.

Estimate, and a caveat that must be stated plainly. The same two methods give a band of 89.84–89.96 fF, within 0.07% of the anchor on either side. **This band contains the anchor by construction.** The corrected series is the first series rescaled: round by round, the ratio of the two is constant to a relative spread of 8.1×10^{-5} (largest ratio over smallest, minus one; our computation from the two committed series). An extrapolation rule applied to both therefore returns, to about that accuracy, the first run's extrapolated value times that ratio, and the target was set to the anchor times the inverse of the first run's extrapolated bias. The width of the band is the spread between extrapolation methods. It is not an uncertainty. What the second run does establish is the upper bound above, and that the relative discretization bias carries over between two nearby designs to about 10^{-4}. At slower assumed rates the same series extrapolates to 89.394848 fF ($\alpha = 1/3$) and 89.678525 fF ($\alpha = 1/2$), which is 0.57% and 0.25% below the anchor. Only the upper side is proven. A two-sided bound needs a guaranteed lower bound, for instance from a complementary-energy formulation of the same capacitance, or a substantially finer or higher-order solve. Neither exists here; it is future work (issue #1062).

Cross-solver check: the same discretization. Palace, run with second-order elements on the exported adapted mesh of the primary design, gives 89.900294 fF against geode-fem's 89.900518 fF, a relative difference of 2.5×10^{-6}. On the resolution-corrected design it gives 90.979877 against 90.980104 fF, the same relative difference. Palace adds 3,777 elements by local bisection at the interior boundaries, so the two codes share the polygonal geometry and coefficients, not one mesh. This shows that two solvers agree on the same discretization. It is not evidence about the discretization error, which the held-out series measures. A third-order Palace run was killed (signal 9) during its error estimation and produced no complete capacitance matrix; its log is retained and it is not evidence. Two adaptive Palace steps left the island capacitance unchanged while the feedline self-capacitance moved by 1.90%, which suggests the marking concentrated on the feedline; we do not count them as a convergence check either.

What is not claimed. No continuum value is certified for either design. No scalar-to-tensor equivalence, layout regeneration or re-meshing, full-wave optimization, experimental validation, global optimum or speedup against Palace is claimed. The geometry has zero-thickness sheet conductors and a natural zero-charge outer boundary, and "converged" means converged for the

polygonal surfaces of the committed mesh. The choice of P2 elements and ten refinement cycles was made after exploratory P1 runs, which are retained in the artifact directory; it was not preregistered.

Roadmap. Two extensions are planned and not built. The first is the two-sided bound above. The second is a second design target, the island–claw coupling ratio $\beta = C_g/C_\Sigma$ (issue #1037): the coarse-mesh optimizer interface accepts several targets and is unit-tested with two targets and three parameters, but no two-target experiment on the device exists.

10 Roadmap: Differentiating the Eigenmode Branch

The LOM branch above differentiates an *electrostatic* solve. The transmon's dynamical figures-of-merit — resonator frequency, participation-ratio and EPR self/cross-Kerr (Minev et al., 2021a; Nigg et al., 2012) — live in the *eigenmode* problem, and that branch is **not differentiable yet** in geode-fem. The paper commits to the explicit split:

Branch	Observables	Differentiable?
LOM (electrostatic)	E_C, $\alpha \approx -E_C$, coupling C	**yes** (Sections 8–9)
eigenmode / EPR	resonator f, participation, Kerr	**not yet** (this section)

Differentiating an eigenpair needs the adjoint-eigenpair (Hellmann–Feynman / Nelson) formulas (Nelson, 1976), which require solving the interior eigenproblem robustly at the physical shift — and geode-fem's matrix-free interior eigensolve is blocked at exactly that shift. The blocker is measured and committed (`benchmarks/transmon_bench_cpu/sigma4p5_deepshift_characterization.md`): at $\sigma = 4.5$ GHz on the 133k-DOF fixture, a *single* inner AMS-preconditioned MINRES solve (outer Lanczos step 0) never converges — it plateaus on a nearly flat residual tail below $\sim 10^{-5}$ and was still running at 14,300 inner iterations ($\|r\| \approx 2.6 \times 10^{-6}$) when killed. The plateau is *coarse-solve-invariant* — an exact coarse solve stalls the same way — so the limiter is the SPD-proxy preconditioner $K + |\sigma|M$ itself, not the coarse solve. Nor is there a single inner tolerance that is both cheap and correct: at an inner tolerance of 10^{-2} the entire eigensolve completes in $\sim$56 s but returns the non-physical $\lambda \approx 0$ gradient near-kernel cluster (0.42–0.64 GHz), not the physical band. The named path — not built, cited as method, not result — is a preconditioned Helmholtz-projected Jacobi–Davidson iteration, applying the gradient-kernel projection inside the correction equation, for which Liang et al. (2026) provide the current adaptive-multilevel treatment of exactly this singular Maxwell eigenproblem; combined with the Nelson/Hellmann–Feynman eigenpair derivatives, that is the route to differentiable eigenmode/EPR quantities. Until it lands, no claim in this paper extends to derivatives of the eigenmode spectrum.

11 Performance Context

The timings below qualify the cost of the capability. Eigenmode, driven CPU and driven GPU results are separate experiments; none measures the optimizer's cost or establishes a general solver ranking.

11.1 CPU eigenmodes: same box, verified mode content

The earlier timing request, six modes at 4.5 GHz, did not establish equivalent work. In the mode-logged re-run it returns four near-kernel modes, the 3.45 GHz qubit mode and only one of Palace's six modes in geode-fem, whereas Palace, which searches only above its target, returns its six. The historical run has no committed mode log or CPU-time record; its mode content and actual thread

Table 6: Mode-verified CPU timings on one m6i.4xlarge: median wall seconds / measured core-seconds. All rows contain the six reference modes of Table 1. The solvers start from the same input mesh but Palace solves 183,891 unknowns and geode-fem 156,863; the timings are not normalized for that. Three repeats per cell except Palace one-rank Save=0 (one run). The two oracle-tuned requests are marked †; the historical port-aware route is the pre-port-subspace implementation. Palace always requests six modes above 4.5 GHz.

Request	Extra modes	1 wide	8 wide
geode 30 at 4.5 GHz†	24	34.45 / 34.44	29.08 / 40.71
geode 6 at 20 GHz†	0	33.21 / 33.20	27.26 / 38.87
geode port-aware 8 at 4.5 GHz	2	67.26 / 67.25	51.12 / 91.72
Palace, Save=6	0	186.84 / 186.82	33.35 / 266.66
Palace, Save=0	0	113.00 / 112.98	23.64 / 188.89

count cannot be recovered. We therefore withdraw the earlier claim that one geode core beats eight Palace ranks with roughly twelve times fewer core-seconds. This does not change the separate frequency agreement credential of Section 5.

The replacement experiment ran on one otherwise idle EC2 `m6i.4xlarge` in `us-east-1` on 2026-10-08, with geode `61e8571e` and Palace `fba6a5b`. One-wide runs were OS-pinned to one CPU; eight-wide runs used one hardware thread per physical core. Wall clock covers the whole solver process; core-seconds are measured user plus system CPU time, not wall time multiplied by requested width. Table 6 and Figure 7 retain all three geode requests that recover the six reference modes. Two are tuned using the reference spectrum: requesting thirty modes at 4.5 GHz, or moving the shift to 20 GHz. The older port-aware extraction route does not use reference frequencies to choose its request, but returns two further modes alongside the six (one near-kernel mode and the qubit mode); matching the six still uses Palace. Palace's target is a lower bound, whereas geode's shift is a centre, so identical numerical requests need not select identical spectra.

Two asymmetries qualify every comparison in Table 6. First, the solvers do not solve the same number of unknowns. Palace refines the input mesh around the metal sheets, and every timed Palace run assembles 183,891 Nédélec unknowns against geode-fem's 156,863 (133,108 interior), about 17% more (Section 4.1). The table does not correct for this, and the difference favours geode-fem. Second, the reference set is Palace's six modes above 4.5 GHz. That set includes the 17.49 GHz port-sheet mode and excludes the qubit mode, which Palace cannot return at this target (Section 6). The "Extra modes" column counts returned modes outside that set. In the thirty-mode and port-aware rows one of the extras is the qubit mode, which the committed artifact still labels `spurious_port`; the others are near-kernel modes.

At eight wide Palace without field output takes 23.64 s, less than any of these geode routes (27.26–51.12 s). With field output, Palace takes 33.35 s and the two oracle-tuned routes are faster. At one wide the geode routes take less wall and CPU time, but neither comparison isolates a solver kernel: geode writes no fields and has no counterpart to Palace's error estimator, which remains enabled even with Save=0 (about 40.2 s at one rank and 5.0 s at eight). MPI busy-waiting also affects measured Palace core-seconds. Palace also carries about 17% more unknowns in every cell. No single route is selected as a headline winner.

The Palace build uses SuperLU-preconditioned GMRES, not AMS: the pinned source and build configuration resolve `Default` to SuperLU. A later, separate m6i session with a tuned Palace build found only a 1.01–1.05× same-box wall-time improvement over its untuned control; cross-session differences must not be credited to compiler tuning. Those runs also predate the new port-subspace solver. The local update in Section 6.2 cannot be substituted into this same-box timing table.

At 1,157,564 DOFs, a historical geode direct run completed on a 128 GB box in 565.5 s with

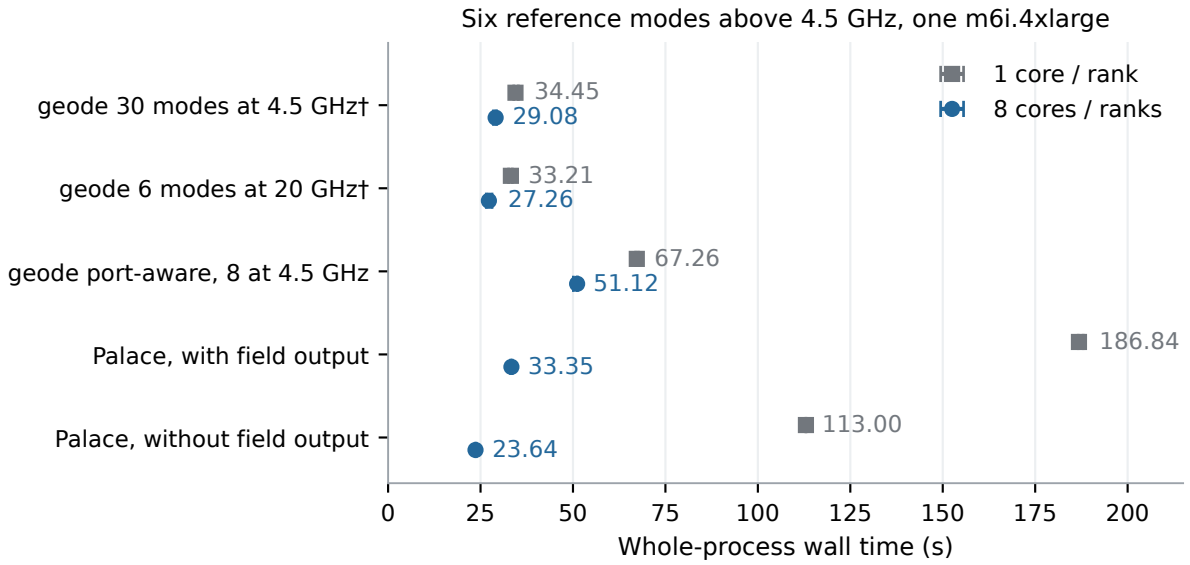


Figure 7: Whole-process wall time on one m6i.4xlarge for the three mode-recovering geode requests and for Palace with and without field output (Table 6), on one linear axis for one and eight cores or ranks. Markers are medians; whiskers span the observed minimum and maximum. Every request returns the six reference modes of Table 1, which include the 17.49 GHz port-sheet artifact and exclude the qubit mode. The two oracle-tuned requests (marked †) are not a reference-free recipe for an unknown device. Save=0 removes Palace's field output but retains its error estimator, and Palace solves about 17% more unknowns from the same input mesh, so these are workload comparisons, not solver-kernel speedups.

92.2 GB peak RSS. Palace's earlier eight-rank run recorded 423.12 s and about 33 GB aggregate memory. These are different-host observations, and the geode request lacks mode-content verification. The DOF count is geode-fem's; we infer that Palace's internal refinement applies to that mesh as well, and its unknown count for that run is not quoted here. They establish the direct factorization's large memory requirement, not a controlled performance crossover or equivalent physical-mode throughput.

11.2 CPU driven solve: AMS at scale

A separate f64 experiment uses assembled COCG on the lossy driven cube at $\omega = 0.10$, on one dedicated `r6i.4xlarge`, with geode `64d44269`. AMS reduces one-core linear-solve time relative to Jacobi by 3.4–4.2× across 25,695–897,650 edges. It reduces iterations from 2969–13662 to 39–105, but does not eliminate iteration growth. Every geode solve meets an explicitly recomputed relative-residual tolerance of 10^{-8}.

At 462,520 edges AMS needs about 3.4 GB versus Jacobi's 1.05 GB. At 897,650 edges it needs 8.5 GB versus 2.0 GB, and its setup alone takes 53.488 s. Eight threads reduce setup but do not accelerate geode's Krylov iteration; at 462,520 edges its AMS solve takes 50.4 s versus Palace AMS's 4.2 s on eight ranks. The one-core geode AMS advantage at the smallest size does not persist at scale: Palace AMS takes 53.3 s against 181.3 s at the largest size. These measurements concern one structured first-order fixture, assembled CPU operators and low frequencies. They establish neither GPU performance nor a remedy for the interior eigenproblem.

Table 7: Driven CPU-f64 linear-solve seconds on one r6i.4xlarge, one core/rank, median of two geode and three Palace runs, except 897,650 edges (one run each). geode includes operator assembly, preconditioner setup, iteration and residual check; Palace sums its Setup, Preconditioner and Linear Solve timers, excluding Operator Construction. These are not whole-process timings.

Edges	geode Jacobi	geode AMS	Palace AMS	Palace SuperLU
25,695	2.4	0.7	1.0	1.4
102,024	22.2	5.3	5.0	8.9
462,520	227.3	56.4	26.4	85.9
897,650	624.3	181.3	53.3	235.8

11.3 GPU cell: correctness proven, scaling an honest negative

geode-fem's GPU path runs the *driven* solve (matrix-free apply + on-device COCG), not the eigensolve of the credential, and all GPU execution is CUDA-f32 (Section 3). On 2026-07-14, geode-fem code executed on a physical GPU for the first time (EC2 `g6e.xlarge`, 1× NVIDIA L40S 46 GB, driver 595.71.05, CUDA 13.2): the two CUDA-f32 correctness smokes — the matrix-free Nédélec matvec and the on-device COCG — pass against the f64 CPU reference within f32-appropriate tolerances (15.0 and 3.3 s wall, durations dominated by GPU init and kernel JIT — correctness evidence, not performance).

The scaling study (committed as `benchmarks/gpu_driven_scaling/results.toml`; a σ-lossy lumped-port cube fixture at fixed frequency, mesh-refined from 1,854 to 25,695 edges, all sixteen cells on the same host) is an honest negative: **GPU-f32 matrix-free is slower than both assembled CPU configurations at every measured size, and slower than CPU matrix-free at every size except the largest**, where the two solve-only medians are at parity (81.764 versus 82.056 s; Table 8). Computed from the artifact's unrounded medians, the assembled-CSR COCG on CPU is 136× faster at 1,854 edges (displayed: 0.032 versus 4.39 s) and still 44× faster at 25,695 (1.86 versus 81.8 s); even the direct LU is 13.5× faster at the top size (6.04 versus 81.8 s). The mechanism is kernel-launch domination (∼28 ms per iteration at the top size versus ∼0.63 ms for assembled-CSR on CPU) — 25.7k edges is far too small to fill an L40S — and f32 accuracy compounds the verdict (relative error degrading to 1.2×10^{-3} with refinement; nondeterministic convergence at the top size). One directional positive survives, and it is the one the architecture predicts: against the *same matrix-free algorithm on CPU*, the GPU's deficit falls monotonically to parity at ∼26k edges (81.8 versus 82.1 s). We state the GPU position as trajectory, not achievement, with tracked levers (kernel fusion; on-device f64/mixed-precision when cubecl ships it; preconditioning) — and we prefer publishing the bound to implying the win. This cell is a cube fixture family, not the transmon eigenproblem, and its CPU baselines (4-vCPU host) are not comparable to the `m6i` numbers of Table 6.

12 Reproducibility

Every artifact behind every result in this paper is committed and scripted.[3] The exceptions are labelled where they occur: the scratch-probe measurements and source readings of Section 6 trace to a committed design note and probe source, not to a regenerable benchmark artifact.

- **Geometry and mesh.** Generated by DeviceLayout.jl v1.15.0's `SingleTransmon` example (AWS Center for Quantum Computing, 2025); the fixture generator is committed, and the

[3]Repository URL and archival DOI for the geode-fem source to be added at submission. TODO(operator).

Table 8: Driven-solve wall-clock scaling, solve-only medians of $n = 3$ timed repetitions (warm-up excluded), same host (`g6e.xlarge`: 4 vCPU + 1× L40S). One fixture (σ-lossy lumped-port cube) at fixed ω, mesh-refined across the columns. f64 iterative configurations at tolerance 10^{-8}; the f32 GPU configuration at 10^{-6} (its recurrence floor sits above 10^{-8}). The accuracy row is the full-field relative L_2 error of the GPU-f32 solution against the Direct-f64 reference. At 25,695 edges the GPU-f32 configuration converged in the committed single-frequency runs but not in its five-point frequency sweep, and run-to-run convergence is nondeterministic at that size (f32 floor). Values verbatim from the committed `benchmarks/gpu_driven_scaling/results.toml`.

	Mesh size (Nédélec edges)			
Configuration	1,854	5,859	13,428	25,695
Direct faer LU (CPU f64), s	0.024	0.203	1.540	6.036
Assembled-CSR COCG (CPU f64), s	0.032	0.198	0.709	1.865
Matrix-free COCG (CPU f64), s	1.652	8.885	30.234	82.056
Matrix-free COCG (GPU f32), s	4.388	13.351	34.381	81.764
GPU-f32 accuracy vs. Direct-f64 (rel. L_2)	1.26×10^{-4}	2.93×10^{-4}	4.60×10^{-4}	1.20×10^{-3}

MSH 4.1 fixture is pinned by SHA-256 (`5b3ff4c3...b33dd`; full hash and provenance in the committed `transmon_smoke.provenance.txt`).

- **The MSH conversion gotcha.** MFEM's MSH 4.1 reader rejects the DeviceLayout multi-entity-block file, so the Palace runs consume an MSH 2.2 conversion (`gmsh -save -format msh2`). The conversion is node-preserving (all 22,684 nodes unchanged, physical groups preserved) and physics-neutral: eigenvalues reproduce bit-for-bit across the conversion. Documented in the committed `palace_config.provenance.txt` because it is exactly the kind of silent pipeline step a same-mesh claim must disclose. A second such step went undisclosed in earlier versions: Palace refines the mesh it reads (4,913 elements added by local bisection, 6,014 duplicated vertices) and solves 183,891 unknowns, as its committed run logs record (`reference/fixtures/transmon_palace/results_p1/palace_run_v22.log`).

- **Credential artifacts.** The Palace configuration and provenance under `reference/fixtures/transmon_palace/`; geode-fem's benchmark definitions, the agreement numbers, the L-doubling tripwire (17.4901 → 12.37 GHz, committed ratio 0.7071), the qubit-mode measurements (3.4528 GHz, $p = 0.9942$; L-doubling ratio 0.7081; stored under the artifact's `[spurious_mode]` key), Palace's `Target` = 2 GHz re-run (`[oracles.palace_target2]`), and the full gauge/projection arc of Section 6.2 (tree–cotree counts and the 1.64% shift; bulk-projection divergence ratio 50.2 and projected norm 1.06×10^{-4}; port-aware acceptance at $\leq 0.029\%$) in `benchmarks/transmon_eigen/results.toml`, each pinned by a release-tier regression test. Palace's raw outputs (`eig.csv`, per-port participation, run log) are committed; a rerun reproduced `eig.csv` bit-for-bit. The re-run's configurations, logs, `eig.csv`, `port-EPR.csv` and provenance are under `reference/fixtures/transmon_palace/results_p1_target2/`. The source reading of Palace's port operator and search window, the port-voltage probe and the planned rank-1 formulation are in `docs/research/transmon-lumped-port-formulation.md`; its probe is scratch source that the build does not compile.

- **Precision of the committed agreement numbers.** `benchmarks/transmon_eigen/results.toml` stores geode-fem frequencies at three decimals while the Palace column carries full precision, so the committed per-mode `rel_err_pct` is computed against the rounded geode-fem values; at full precision the agreement is slightly better (approximately 0.029% worst case). We quote the committed-artifact numbers throughout and disclose the rounding

here. Derived ratios quoted in the text (core-second and speedup multipliers) are computed from the displayed table values unless marked as committed.

- **LOM chain.** The three-conductor capacitance extraction and Koch-exact spectrum on the coarse mesh ($C_\Sigma = 136.7\,\mathrm{fF}$, $E_C = 0.142\,\mathrm{GHz}$, and the boundary and sapphire-model checks of Section 7) in `benchmarks/transmon_quantum/results.toml`; pinned by the release test `crates/geode-core/tests/transmon_quantum.rs`. The mesh-convergence study (P1 to $h/4$, P2 to $h/2$, scalar and tensor permittivity, with its bounds, estimates and memory record) in `benchmarks/transmon_diffopt/c_sigma_convergence.toml`, written by the example `transmon_c_sigma_convergence`; it ran locally in under five minutes; earlier runs of it peaked at 12.8–13.8 GB of resident memory.

- **Adjoint layer.** The four sensitivity modules with their FD-and-mutation test suites: `crates/geode-core/src/adjoint.rs` (PR #573); `crates/geode-core/src/shape.rs` (PRs #575/#586; this module also carries `capacitance_shape_gradient`); `crates/geode-core/src/driven/adjoint.rs` (PR #579); and `crates/geode-core/src/driven/shape.rs` (PR #581).

- **Optimization demonstrations.** The parallel-plate loop proof (full trajectories, closed-form check, fresh-solve confirmation) in `benchmarks/transmon_diffopt/results.toml`. The refined real-device optimization under `benchmarks/transmon_diffopt/refined_optimization/`: each run directory (`p2_cycles10/` and `p2_cycles10_resolution_corrected/`) holds a generated `summary.toml` with the optimizer steps, the final finite-difference checks, the held-out series and every bound and estimate with its rounding direction, plus the Palace configurations, logs and outputs; `validation_summary.json` is a hand-assembled index of the same numbers, and the directory's `README.md` gives the commands (about 20 and 10 minutes locally, with peak resident memory up to 29.4 and 26.0 GB). The coarse-mesh runs: the single-parameter FD validation and bounded anchor attempt in `benchmarks/transmon_diffopt/pad_results.toml`; harmonic motion in `benchmarks/transmon_diffopt/harmonic_results.toml`; the three-parameter gradients in `benchmarks/transmon_diffopt/multiparam_gradient.toml`; and the bounded optimizer runs in `benchmarks/transmon_diffopt/multiparam_results.toml`. All are consumed by `crates/geode-core/tests/transmon_diffopt.rs`, which pins both `summary.toml` files and re-derives their bounds and estimates.

- **Performance artifacts.** The historical CPU cell in `benchmarks/transmon_bench_cpu/results.toml`; mode-verified timing in `benchmarks/transmon_bench_cpu/results_like_for_like_m6i.toml`; tuned-build controls in `benchmarks/transmon_bench_cpu/results_tuned_palace_m6i.toml`; local port-subspace evidence in `benchmarks/transmon_bench_cpu/results_like_for_like_local.toml`; driven CPU AMS results in `benchmarks/gpu_driven_scaling/results_ams_cpu_r6i.toml`; the 1.16M-DOF completion run log at `benchmarks/transmon_bench_cpu/geode_runs_1p16M_2026-07-15.log`; the GPU scaling cell (all sixteen cells, iteration counts, recomputed true residuals, DNF handling, same-host provenance) in `benchmarks/gpu_driven_scaling/results.toml`; the deep-shift eigensolve characterization at `benchmarks/transmon_bench_cpu/sigma4p5_deepshift_characterization.md`.

- **Hardware disclosure.** CPU cell: EC2 `m6i.4xlarge` (8 physical cores / 16 vCPU, 64 GB), `us-east-1` for the mode-verified re-timing (the historical session was in `us-west-2`); driven AMS: `r6i.4xlarge`, `us-east-1`; the 1.16M-DOF completion row: a 128 GB `r6i`-class box (see the committed log). GPU cell: EC2 `g6e.xlarge` (4 vCPU, 30 GB; 1× NVIDIA L40S 46 GB; driver 595.71.05, CUDA 13.2). The Palace `Target` = 2 GHz re-run: a local Apple M3

Ultra, Palace built from the in-repository Docker recipe (arm64), eight MPI ranks; no timing from it is used.

- **Versions.** geode-fem commit `3174015` (headline credential and historical CPU cells; the adjoint/diffopt layers are merged post-`3174015` and cited by PR number); Palace commit `fba6a5b`; DeviceLayout.jl v1.15.0; Burn/burn-cuda 0.21.

13 Discussion

What the credential does and does not establish. In V&V terms (Oberkampf and Roy, 2010; Roache, 1997) the cross-validation is code verification and code-to-code comparison, not validation against nature: two independent implementations agreeing to $\leq 0.033\%$ from the identical input mesh, on discrete problems that differ by Palace's internal refinement, with scaling tripwires on the port operator and on the qubit mode. It does not certify that either solver matches a measured device — that is what Ye et al. (2025)-style experimental benchmarks address — but it rules out the large error class living in formulation, assembly, boundary handling, and eigensolver machinery. It establishes that geode-fem's eigenmode path reproduces a production solver from the same input. It does not establish mesh convergence. In particular the electrostatic observable that this paper differentiates is not converged on this mesh (Section 7), and the eigenmode credential says nothing about that.

What the differentiable-design evidence supports. The proven capability is bounded and we state its boundary. Proven: the 2×2 adjoint matrix at operator level (FD-validated with mutation tests); the capacitance$\to E_C$ chain with its stationarity collapse; a closed Newton loop confirmed by an independent fresh solve on an analytic fixture; a three-parameter optimization of the real device on an adaptively refined P2 discretization, with gradients validated against fully re-solved finite differences at 3.87×10^{-7} or better; and a rigorous upper bound on the converged capacitance of each optimized design (below 89.02 fF, and below 90.08 fF for the resolution-corrected one; Section 9.4). Estimated, not proven: the converged values themselves (about 88.8–88.9 fF and 89.84–89.96 fF). The second band contains the anchor by construction. Not proven, and not claimed: a lower bound on either converged capacitance, and therefore any certified distance to the anchor from below; that the first design is within 1% of the anchor; eigenmode/EPR derivatives (the Section 10 wall); two-target or coupling-constrained optimization; tensor-ε shape chains; a design regenerated by the layout toolchain or re-meshed; and any wall-clock superiority over sweep-based workflows. The capability is also not qubit-specific — shape and material sensitivities of electrostatic and driven-EM observables apply to RF and photonic device design generally; the transmon is the demonstration vehicle with the clearest documented unmet need.

Threats to validity. A common-mode error — both solvers wrong the same way — cannot be excluded by cross-comparison alone, and this paper now contains an example. The 17.49 GHz mode is an artifact of a port representation that both solvers share; the two agree on it to 0.001%, and earlier versions took that agreement, with the f_{LC} anchor, as confirmation of the junction physics. The same versions asserted the absence of a mode that Palace had not been asked for. Agreement with a reference bounds implementation error, not shared modeling choices or the reference's own search window. Every capacitance in this paper is a conforming upper bound, and no lower bound has been computed. The estimates of converged values rest on extrapolations whose rates are not established; at slower assumed rates the resolution-corrected design lands 0.25–0.57% below the anchor. That design's estimate band contains the anchor by construction. The two refined runs share

a mesh family, a marking rule and an element order, so neither confirms the other independently, and Palace's agreement on the adapted mesh is agreement on the same discretization. The choice of P2 elements and ten refinement cycles was made after exploratory runs and was not preregistered. No two-target or coupling-constrained optimization is demonstrated. The parallel-plate tier's one-step convergence is an affine-map property, disclosed as such. The real-device chain uses the scalar trace-averaged sapphire, and the scalar-versus-tensor difference is not converged: it changes sign under refinement (Section 7). Repeat counts and timing boundaries differ across experiments and are stated in each table. The large-scale eigenmode completion used a different 128 GB box and lacks mode-content verification; it cannot establish equivalent-work throughput.

Limitations. (i) One geometry. The eigenmode results are on one mesh at first order; higher-order and refined eigensolves are future work (a $p = 2$ Palace configuration is committed for that purpose). The electrostatic branch has a convergence study, and it shows the committed mesh far from converged for C_Σ (Section 7). (ii) The lossless model leaves readout ports open; resistive terminations (hence complex-symmetric pencils and mode Qs) are out of scope. (iii) The credential eigenpath remains the ungauged committed benchmark, and mode classification still uses the Palace reference spectrum. The shipped port term retains the 17.49 GHz port-sheet artifact; the rank-1 port-voltage formulation that removes it in a probe is planned and not built. The committed artifacts still label the qubit mode as spurious. (iv) The CPU evidence depends on request selection, field output, error estimation and implementation version, and Palace solves about 17% more unknowns from the same input mesh; it supports no unconditional per-core or scaling advantage. (v) GPU evidence covers driven-solve correctness and an honest scaling negative in f32 on a cube fixture family; no GPU eigensolve path exists. (vi) The differentiable branch is electrostatic/LOM only (eigenmode/EPR is roadmap) and scalar-ε for the shape chain. Its real-device result carries upper bounds and extrapolated estimates only, for the polygonal geometry of the committed mesh and with the junction's 5.5 fF outside C_Σ. The refined optimization runs in native f64 with hand-derived element integrals, outside the Burn tape.

14 Conclusion

The discrete-adjoint layer supplies solver-derived material and geometry gradients for electrostatic and driven FEM operators, and the stationary capacitance derivative closes the geometry-to-charging-energy loop. On the real device, a finite-difference-validated P2 shape gradient on an adaptively refined mesh reaches the 89.9 fF anchor in three accepted steps. That result is for the scalar-permittivity electrostatic C_Σ without the junction's 5.5 fF, and its gradient is a hand-derived equilibrium-energy shape derivative solved iteratively in native f64, outside the Burn tape and without a factorization. Held-out refinement then shows what that design is worth: its converged capacitance is rigorously below 89.02 fF and is estimated at about 88.8–88.9 fF. A resolution-corrected design converges below 90.08 fF, with an estimate band that contains the anchor by construction. Only upper bounds are proven; no continuum value is certified, and a two-sided bound is the next experiment. The convergence study behind these statements is a result of its own: the committed mesh overstates the scalar-permittivity C_Σ by at least 54.69 fF and the tensor-permittivity one by at least 47.18 fF, the converged device lies below the anchor and not above it in both models (below 83.02 fF and below 89.50 fF), and the coarse-mesh optimizations that lowered C_Σ to reach the anchor moved the wrong way while passing every self-consistency check. Same-input-mesh agreement with Palace supplies a forward-solver credential that now includes the transmon qubit mode, which earlier versions wrongly reported as absent. The updated mode-verified timing and driven AMS experiments delimit the computational cost without claiming a general performance

advantage. Eigenmode/EPR derivatives remain future work.

Acknowledgments. geode-fem's development was AI-orchestrated (agent-driven implementation with human direction and review); this benchmark program and paper follow the same discipline: results are traced to committed, human-auditable artifacts, and statements that rest on a scratch probe, on a reading of Palace's source, or on inference are labelled as such.

A The historical optimization figure, and figure provenance

Figure 8 keeps the two earliest optimization tiers in view: the analytic loop proof of Section 9.1 and the single-parameter runs on the coarse device mesh of Section 9.2. Panel (b) is historical. Its anchor shortfall is an artifact of the unconverged coarse mesh (Section 7), and the mesh-validity budget it shows belongs to that fixed-topology map on that mesh; it is not a limit on the refined optimization of Figure 6.

Every plotted data value in this paper is read by the figure scripts in `figures/src/` from a committed artifact, with one exception: the optimized island outline of Figure 6c–e. It comes from the committed mesh moved to the primary design, as exported by the refined run; the artifact directory does not commit that mesh (its meshes are regenerated by the example). Its 139 island triangles are kept with the figure sources as a snapshot that records the SHA-256 of the exported mesh, and the figure verifier checks that the snapshot's extent reproduces the committed island length and pad width exactly. The scripts need Python 3.11 or newer, NumPy and Matplotlib, and no private plotting style.

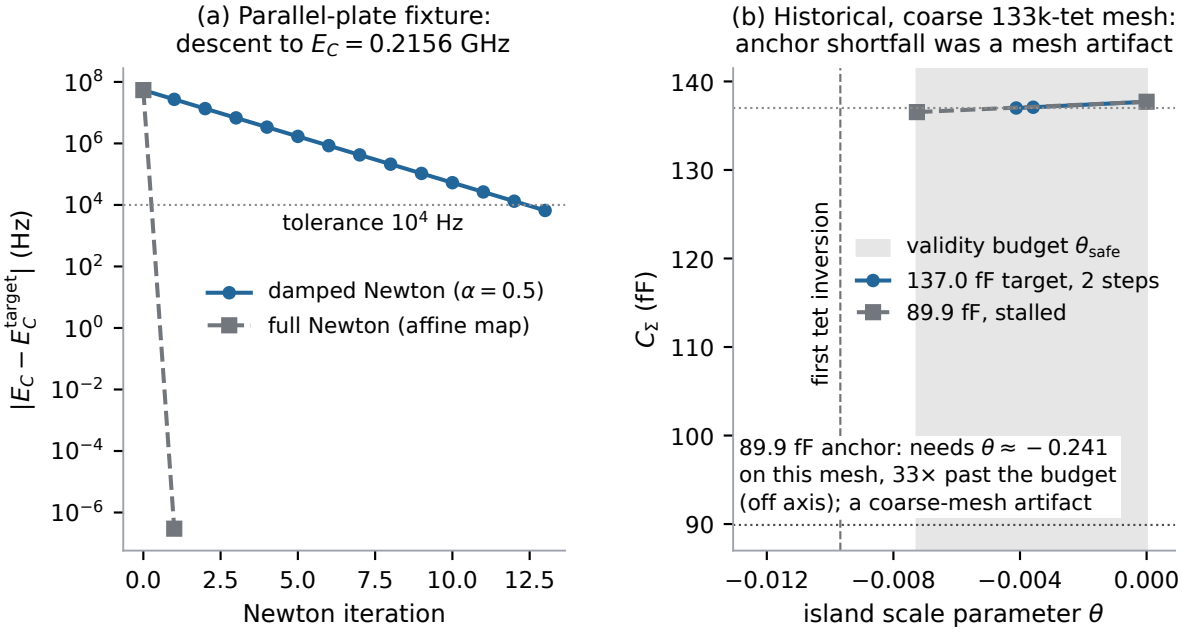


Figure 8: Gradient-based optimization to a target E_C via the discrete-adjoint shape gradient. (a) Parallel-plate loop proof: $|E_C - E_C^{\text{target}}|$ versus Newton iteration; full Newton converges in one step (the parametrization is affine — expected, and disclosed), the damped $\alpha = 0.5$ run supplies the 13-step descent curve; the converged design is confirmed by an independent fresh forward solve at relative error 1.4×10^{-15} (`benchmarks/transmon_diffopt/results.toml`). (b) Historical, on the coarse 133k-tet DeviceLayout mesh under the fixed-topology island-scale parameter θ: a within-budget demonstration target ($C_\Sigma = 137.0\,\text{fF}$) converges in two genuine (non-affine) Newton steps, fresh-solve confirmed at 5.6×10^{-6}; the bounded run toward the 89.9 fF design anchor stalls honestly at the mesh validity budget $\theta_{\text{safe}} = -0.0073$ (first tetrahedron inversion at $\theta = -0.0097$), 33× short of the $\theta \approx -0.241$ the anchor requires on that mesh (`benchmarks/transmon_diffopt/pad_results.toml`). The shortfall is an artifact of the unconverged mesh: converged, the device is already below the anchor (Section 7), and Section 9.4 reports the refined optimization.

References

R. Albanese and G. Rubinacci. Solution of three dimensional eddy current problems by integral and differential methods. *IEEE Transactions on Magnetics*, 24(1):98–101, 1988. doi: 10.1109/20.43865. URL https://doi.org/10.1109/20.43865.

Martin S. Alnæs, Anders Logg, Kristian B. Ølgaard, Marie E. Rognes, and Garth N. Wells. Unified form language. *ACM Transactions on Mathematical Software*, 40(2):1–37, 2014. doi: 10.1145/2566630. URL https://doi.org/10.1145/2566630.

Robert Anderson, Julian Andrej, Andrew Barker, Jamie Bramwell, Jean-Sylvain Camier, Jakub Cerveny, Veselin Dobrev, Yohann Dudouit, Aaron Fisher, Tzanio Kolev, Will Pazner, Mark Stowell, Vladimir Tomov, Ido Akkerman, Johann Dahm, David Medina, and Stefano Zampini. MFEM: a modular finite element methods library. *Computers & Mathematics with Applications*, 81:42–74, 2021. doi: 10.1016/j.camwa.2020.06.009.

Julian Andrej, Nabil Atallah, Jan-Phillip Bäcker, Jean-Sylvain Camier, Dylan Copeland, Veselin Dobrev, Yohann Dudouit, Tobias Duswald, Brendan Keith, Dohyun Kim, Tzanio Kolev, Boyan Lazarov, Ketan Mittal, Will Pazner, Socratis Petrides, Syun'ichi Shiraiwa, Mark Stowell, and Vladimir Tomov. High-performance finite elements with mfem. *The International Journal of High Performance Computing Applications*, 2024. doi: 10.1177/10943420241261981. URL https://doi.org/10.1177/10943420241261981.

Alexander Arsenovic, Julien Hillairet, Jackson Anderson, Henrik Forsten, Vincent Ries, Michael Eller, Noah Sauber, Robert Weikle, William Barnhart, and Franz Forstmayr. scikit-rf: An open source python package for microwave network creation, analysis, and calibration [speaker's corner]. *IEEE Microwave Magazine*, 23(1):98–105, 2022. doi: 10.1109/mmm.2021.3117139. URL https://doi.org/10.1109/mmm.2021.3117139.

AWS Center for Quantum Computing. DeviceLayout.jl: computer-aided design of quantum integrated circuits. https://github.com/aws-cqc/DeviceLayout.jl, 2025. v1.15.0 used for geometry/mesh generation in this work. No journal or arXiv artifact exists as of 2026-07-14 (litsearch absence check).

AWS Quantum Technologies Blog. AWS releases open-source software Palace for cloud-based electromagnetics simulations of quantum computing hardware. https://aws.amazon.com/blogs/quantum-computing/aws-releases-open-source-software-palace-for-cloud-based-electromagnetics-simulations-of-quantum-computing-hardware/, 2023.

Jed Brown, Ahmad Abdelfattah, Valeria Barra, Natalie Beams, Jean-Sylvain Camier, Veselin Dobrev, Yohann Dudouit, Leila Ghaffari, Tzanio Kolev, David Medina, Will Pazner, Thilina Ratnayaka, Jeremy Thompson, and Stan Tomov. libceed: Fast algebra for high-order element-based discretizations. *Journal of Open Source Software*, 6(63):2945, 2021. doi: 10.21105/joss.02945. URL https://doi.org/10.21105/joss.02945.

Tianqi Chen, Thierry Moreau, Ziheng Jiang, Lianmin Zheng, Eddie Yan, Meghan Cowan, Haichen Shen, Leyuan Wang, Yuwei Hu, Luis Ceze, Carlos Guestrin, and Arvind Krishnamurthy. Tvm: An automated end-to-end optimizing compiler for deep learning, 2018. URL https://arxiv.org/abs/1802.04799.

Mingyuan Chi and Shizheng Wen. torch-sla: Differentiable sparse linear algebra with adjoint solvers and sparse tensor parallelism for pytorch, 2026. URL https://arxiv.org/abs/2601.13994.

P. Dular, C. Geuzaine, F. Henrotte, and W. Legros. A general environment for the treatment of discrete problems and its application to the finite element method. *IEEE Transactions on Magnetics*, 34(5):3395–3398, 1998. doi: 10.1109/20.717799. URL https://doi.org/10.1109/20.717799.

Axel M. Eriksson, Lukas J. Splitthoff, Harsh Vardhan Upadhyay, Pietro Campana, Niranjan Pittan Narendiran, Kunal Helambe, Linus Andersson, and Simone Gasparinetti. Automated, physics-guided, multi-parameter design optimization for superconducting quantum devices, 2025. URL https://arxiv.org/abs/2508.18027.

Christophe Geuzaine and Jean-François Remacle. Gmsh: a 3-D finite element mesh generator with built-in pre- and post-processing facilities. *International Journal for Numerical Methods in Engineering*, 79(11):1309–1331, 2009. doi: 10.1002/nme.2579.

Alec M. Hammond, Ardavan Oskooi, Mo Chen, Zin Lin, Steven G. Johnson, and Stephen E. Ralph. High-performance hybrid time/frequency-domain topology optimization for large-scale photonics inverse design. *Optics Express*, 30(3):4467, 2022. doi: 10.1364/oe.442074. URL https://doi.org/10.1364/oe.442074.

Vicente Hernandez, Jose E. Roman, and Vicente Vidal. SLEPc: a scalable and flexible toolkit for the solution of eigenvalue problems. *ACM Transactions on Mathematical Software*, 31(3):351–362, 2005. doi: 10.1145/1089014.1089019.

Ralf Hiptmair and Jinchao Xu. Nodal auxiliary space preconditioning in H(curl) and H(div) spaces. *SIAM Journal on Numerical Analysis*, 45(6):2483–2509, 2007. doi: 10.1137/060660588.

Johannes Hoffmann, Christian Hafner, Patrick Leidenberger, Jan Hesselbarth, and Sven Burger. Comparison of electromagnetic field solvers for the 3d analysis of plasmonic nano antennas, 2009. URL https://arxiv.org/abs/0907.3570.

Philipp Holl, Vladlen Koltun, and Nils Thuerey. Learning to control pdes with differentiable physics, 2020. URL https://arxiv.org/abs/2001.07457.

Miklós Homolya, Lawrence Mitchell, Fabio Luporini, and David A. Ham. Tsfc: A structure-preserving form compiler. *SIAM Journal on Scientific Computing*, 40(3):C401–C428, 2018. doi: 10.1137/17m1130642. URL https://doi.org/10.1137/17m1130642.

Yuanming Hu, Luke Anderson, Tzu-Mao Li, Qi Sun, Nathan Carr, Jonathan Ragan-Kelley, and Frédo Durand. Difftaichi: Differentiable programming for physical simulation, 2019. URL https://arxiv.org/abs/1910.00935.

Tyler W. Hughes, Ian A. D. Williamson, Momchil Minkov, and Shanhui Fan. Forward-mode differentiation of maxwell's equations. *ACS Photonics*, 6(11):3010–3016, 2019. doi: 10.1021/acsphotonics.9b01238. URL https://doi.org/10.1021/acsphotonics.9b01238.

Jiantao Jiang, Yanli Wang, Yifan Wang, and Hehu Xie. Fieldtnn-based machine learning method for maxwell eigenvalue problems, 2024. URL https://arxiv.org/abs/2411.15828.

Sarah Quiñones El Kazdadi. faer: A linear algebra library for the rust programming language. *Journal of Open Source Software*, 11(123):6099, 2026. doi: 10.21105/joss.06099. URL https://doi.org/10.21105/joss.06099.

Yongha Kim, Anthony W. Jung, Sanmun Kim, Kevin Octavian, Doyoung Heo, Chaejin Park, Jeongmin Shin, Sunghyun Nam, Chanhyung Park, Juho Park, Sangjun Han, Jinmyoung Lee, Seolho Kim, Min Seok Jang, and Chan Y. Park. Meent: Differentiable electromagnetic simulator for machine learning, 2024. URL https://arxiv.org/abs/2406.12904.

Robert C. Kirby and Anders Logg. A compiler for variational forms. *ACM Transactions on Mathematical Software*, 32(3):417–444, 2006. doi: 10.1145/1163641.1163644. URL https://doi.org/10.1145/1163641.1163644.

Jens Koch, Terri M. Yu, Jay Gambetta, A. A. Houck, D. I. Schuster, J. Majer, Alexandre Blais, M. H. Devoret, S. M. Girvin, and R. J. Schoelkopf. Charge-insensitive qubit design derived from the Cooper pair box. *Physical Review A*, 76:042319, 2007. doi: 10.1103/PhysRevA.76.042319.

Tzanio V. Kolev and Panayot S. Vassilevski. Parallel auxiliary space AMG for $H(\mathrm{curl})$ problems. *Journal of Computational Mathematics*, 27(5):604–623, 2009. doi: 10.4208/jcm.2009.27.5.013. URL https://doi.org/10.4208/jcm.2009.27.5.013. Crossref records the online-publication year 2018; print year 2009 used here (litsearch cosmetic fix, see version-dir refs.bib header).

Richard B. Lehoucq, Danny C. Sorensen, and Chao Yang. *ARPACK Users' Guide: Solution of Large-Scale Eigenvalue Problems with Implicitly Restarted Arnoldi Methods*. SIAM, 1998. doi: 10.1137/1.9780898719628.

Liang Liang, Minliang Liu, John Elefteriades, and Wei Sun. Pytorch-fea: Autograd-enabled finite element analysis methods with applications for biomechanical analysis of human aorta. *Computer Methods and Programs in Biomedicine*, 238:107616, 2023. doi: 10.1016/j.cmpb.2023.107616. URL https://doi.org/10.1016/j.cmpb.2023.107616.

Qigang Liang, Xuejun Xu, and Qingquan Zhang. Adaptive multilevel methods for the maxwell eigenvalue problem, 2026. URL https://arxiv.org/abs/2603.29718.

Anders Logg and Garth N. Wells. Dolfin: Automated finite element computing. *ACM Transactions on Mathematical Software*, 37(2):1–28, 2010. doi: 10.1145/1731022.1731030. URL https://doi.org/10.1145/1731022.1731030.

Yannik Mahlau, Frederik Schubert, Lukas Berg, and Bodo Rosenhahn. Fdtdx: High-performance open-source fdtd simulation with automatic differentiation. *Journal of Open Source Software*, 11 (117):8912, 2026. doi: 10.21105/joss.08912. URL https://doi.org/10.21105/joss.08912.

J.B. Manges and Z.J. Cendes. A generalized tree-cotree gauge for magnetic field computation. *IEEE Transactions on Magnetics*, 31(3):1342–1347, 1995. doi: 10.1109/20.376275. URL https://doi.org/10.1109/20.376275.

Huan-Ting Meng, Bao-Lin Nie, Steven Wong, Charles Macon, and Jian-Ming Jin. Gpu accelerated finite-element computation for electromagnetic analysis. *IEEE Antennas and Propagation Magazine*, 56(2):39–62, 2014. doi: 10.1109/map.2014.6837065. URL https://doi.org/10.1109/map.2014.6837065.

Zlatko K. Minev, Zaki Leghtas, Shantanu O. Mundhada, Lysander Christakis, Ioan M. Pop, and Michel H. Devoret. Energy-participation quantization of Josephson circuits. *npj Quantum Information*, 7:131, 2021a. doi: 10.1038/s41534-021-00461-8.

Zlatko K. Minev, Thomas G. McConkey, Maika Takita, Antonio D. Corcoles, and Jay M. Gambetta. Circuit quantum electrodynamics (cqed) with modular quasi-lumped models, 2021b. URL https://arxiv.org/abs/2103.10344.

Sean Molesky, Zin Lin, Alexander Y. Piggott, Weiliang Jin, Jelena Vucković, and Alejandro W. Rodriguez. Inverse design in nanophotonics. *Nature Photonics*, 12(11):659–670, 2018. doi: 10.1038/s41566-018-0246-9. URL https://doi.org/10.1038/s41566-018-0246-9.

Jean-Claude Nédélec. Mixed finite elements in $\mathbb{R}^3$. *Numerische Mathematik*, 35(3):315–341, 1980. doi: 10.1007/BF01396415.

Richard B. Nelson. Simplified calculation of eigenvector derivatives. *AIAA Journal*, 14(9):1201–1205, 1976. doi: 10.2514/3.7211. URL https://doi.org/10.2514/3.7211.

Simon E. Nigg, Hanhee Paik, Brian Vlastakis, Gerhard Kirchmair, S. Shankar, Luigi Frunzio, M. H. Devoret, R. J. Schoelkopf, and S. M. Girvin. Black-box superconducting circuit quantization. *Physical Review Letters*, 108:240502, 2012. doi: 10.1103/PhysRevLett.108.240502.

William L. Oberkampf and Christopher J. Roy. *Verification and Validation in Scientific Computing*. Cambridge University Press, 2010. doi: 10.1017/cbo9780511760396. URL https://doi.org/10.1017/cbo9780511760396.

Palace developers, AWS Center for Quantum Computing. Palace: PArallel LArge-scale Computational Electromagnetics. https://github.com/awslabs/palace, 2023. Open-source 3D finite-element solver; v0.13-class, git fba6a5b used in this work. No associated journal article exists as of 2026-07 (absence re-confirmed by litsearch, 2026-07-14).

Greg Peairs. Schematic-driven design of a quantum processor with DeviceLayout.jl. JuliaCon Global 2025, https://pretalx.com/juliacon-2025/talk/7EYFPN/, 2025. Recording: https://www.youtube.com/watch?v=-LohBVtFdwU.

Greg Peairs and Hugh Carson. Design quantum integrated circuits with open-source software DeviceLayout.jl from AWS. https://aws.amazon.com/blogs/quantum-computing/design-quantum-integrated-circuits-with-open-source-software-devicelayout-jl-from-aws/, 2025. AWS Quantum Technologies Blog; source of the transmon+resonator geometry and the reference workflow reproduced here.

Sofia Ponomareva, Adelin Patoux, Clément Majorel, Antoine Azéma, Aurélien Cuche, Christian Girard, Arnaud Arbouet, and Peter R. Wiecha. Torchgdm: A gpu-accelerated python toolkit for multi-scale electromagnetic scattering with automatic differentiation, 2025. URL https://arxiv.org/abs/2505.09545.

Jonathan Ragan-Kelley, Connelly Barnes, Andrew Adams, Sylvain Paris, Frédo Durand, and Saman Amarasinghe. Halide. *Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design and Implementation*, pages 519–530, 2013. doi: 10.1145/2491956.2462176. URL https://doi.org/10.1145/2491956.2462176.

Taha Rajabzadeh, Zhaoyou Wang, Nathan Lee, Takuma Makihara, Yudan Guo, and Amir H. Safavi-Naeini. Analysis of arbitrary superconducting quantum circuits accompanied by a python package: Sqcircuit. *Quantum*, 7:1118, 2023. doi: 10.22331/q-2023-09-25-1118. URL https://doi.org/10.22331/q-2023-09-25-1118.

Taha Rajabzadeh, Alex Boulton-McKeehan, Sam Bonkowsky, David I. Schuster, and Amir H. Safavi-Naeini. A general framework for gradient-based optimization of superconducting quantum circuits using qubit discovery as a case study, 2024. URL https://arxiv.org/abs/2408.12704.

Florian Rathgeber, David A. Ham, Lawrence Mitchell, Michael Lange, Fabio Luporini, Andrew T. T. Mcrae, Gheorghe-Teodor Bercea, Graham R. Markall, and Paul H. J. Kelly. Firedrake. *ACM Transactions on Mathematical Software*, 43(3):1–27, 2017. doi: 10.1145/2998441. URL https://doi.org/10.1145/2998441.

P. J. Roache. Quantification of uncertainty in computational fluid dynamics. *Annual Review of Fluid Mechanics*, 29(1):123–160, 1997. doi: 10.1146/annurev.fluid.29.1.123. URL https://doi.org/10.1146/annurev.fluid.29.1.123.

Martin F. Schubert, Alfred K. C. Cheung, Ian A. D. Williamson, Aleksandra Spyra, and David H. Alexander. Inverse design of photonic devices with strict foundry fabrication constraints, 2022. URL https://arxiv.org/abs/2201.12965.

Sadman Shanto, Andre Kuo, Clark Miyamoto, Haimeng Zhang, Vivek Maurya, Evangelos Vlachos, Malida Hecht, Chung Wa Shum, and Eli Levenson-Falk. Squadds: A validated design database and simulation workflow for superconducting qubit design. *Quantum*, 8:1465, 2024. doi: 10.22331/q-2024-09-09-1465. URL https://doi.org/10.22331/q-2024-09-09-1465.

David Sommers, Zach Degnan, Divita Gautam, Yi-Hsun Chen, Chun-Ching Chiu, Arkady Fedorov, and Prasanna Pakkiam. Open-source highly parallel electromagnetic simulations for superconducting circuits, 2025. URL https://arxiv.org/abs/2511.01220.

Tracel AI. Burn: a comprehensive deep-learning and tensor-computation framework in Rust. https://github.com/tracel-ai/burn, 2024. v0.21.0 used. No journal/arXiv paper exists (litsearch check 2026-07-14): the repo CITATION.cff lists authors Simard, Fortier-Dubois, Tadjibaev, Lagrange + Burn Framework Contributors, with no DOI and no preferred-citation paper — cited as software.

Shizheng Wen, Mingyuan Chi, Tianwei Yu, Ben Moseley, Mike Yan Michelis, Pu Ren, Hao Sun, and Siddhartha Mishra. Learning, solving and optimizing pdes with tensorgalerkin: an efficient high-performance galerkin assembly algorithm, 2026. URL https://arxiv.org/abs/2602.05052.

Tianju Xue, Shuheng Liao, Zhengtao Gan, Chanwook Park, Xiaoyu Xie, Wing Kam Liu, and Jian Cao. Jax-fem: A differentiable gpu-accelerated 3d finite element solver for automatic inverse design and mechanistic data science. *Computer Physics Communications*, 291:108802, 2023. doi: 10.1016/j.cpc.2023.108802. URL https://doi.org/10.1016/j.cpc.2023.108802.

Jiale Ye, Jiaheng Wang, and Yu-xi Liu. Electromagnetic feature extraction in superconducting quantum circuits: An open-source finite-element workflow using palace, 2025. URL https://arxiv.org/abs/2511.09041.